

text-expertise-model

February 17, 2025

```
[1]: from sklearn.model_selection import train_test_split
import pandas as pd
import numpy as np
!pip install nltk
!pip install nltk matplotlib
train_df = pd.read_csv("train.csv", on_bad_lines='skip') # Skip bad lines
val_df = pd.read_csv("val.csv", on_bad_lines='skip') # Skip bad lines
test_df = pd.read_csv("test.csv", on_bad_lines='skip')
merged_df = pd.concat([train_df, val_df, test_df], ignore_index=True)
```

Requirement already satisfied: nltk in /usr/local/lib/python3.11/dist-packages (3.9.1)

Requirement already satisfied: click in /usr/local/lib/python3.11/dist-packages (from nltk) (8.1.8)

Requirement already satisfied: joblib in /usr/local/lib/python3.11/dist-packages (from nltk) (1.4.2)

Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.11/dist-packages (from nltk) (2024.11.6)

Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from nltk) (4.67.1)

Requirement already satisfied: nltk in /usr/local/lib/python3.11/dist-packages (3.9.1)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (3.10.0)

Requirement already satisfied: click in /usr/local/lib/python3.11/dist-packages (from nltk) (8.1.8)

Requirement already satisfied: joblib in /usr/local/lib/python3.11/dist-packages (from nltk) (1.4.2)

Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.11/dist-packages (from nltk) (2024.11.6)

Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from nltk) (4.67.1)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.1)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (4.56.0)

Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
 Requirement already satisfied: numpy>=1.23 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.26.4)
 Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2)
 Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (11.1.0)
 Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.1)
 Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (2.8.2)
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib) (1.17.0)

Frist method : BERT pre-trained

```
[ ]: # Install CUDA Toolkit
!wget https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1804/
  ↪x86_64/cuda-11-0_11.0.3-1_amd64.deb
!dpkg -i cuda-11-0_11.0.3-1_amd64.deb
!apt-key adv --fetch-keys https://developer.download.nvidia.com/compute/cuda/
  ↪repos/ubuntu1804/x86_64/7fa2af80.pub
!apt-get update
!apt-get install cuda

# Install cuDNN (adjust version accordingly)
!wget https://developer.download.nvidia.com/compute/redist/cudnn/v8.0.5/
  ↪cudnn-11.0-linux-x64-v8.0.5.39.tgz
!tar -xzf cudnn-11.0-linux-x64-v8.0.5.39.tgz
!cp cuda/include/cudnn*.h /usr/local/cuda/include
!cp cuda/lib64/libcudnn* /usr/local/cuda/lib64
!chmod a+r /usr/local/cuda/include/cudnn*.h /usr/local/cuda/lib64/libcudnn*

!nvidia-smi

import torch
torch.cuda.is_available()
```

```
--2025-01-20 03:52:26-- https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1804/x86_64/cuda-11-0_11.0.3-1_amd64.deb
Resolving developer.download.nvidia.com (developer.download.nvidia.com)...
23.215.7.4, 23.215.7.28
Connecting to developer.download.nvidia.com
(developer.download.nvidia.com)|23.215.7.4|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 2446 (2.4K) [application/x-deb]
Saving to: 'cuda-11-0_11.0.3-1_amd64.deb'
```

cuda-11-0_11.0.3-1_ 100%[=====>] 2.39K --.-KB/s in 0s

2025-01-20 03:52:27 (1.23 GB/s) - 'cuda-11-0_11.0.3-1_amd64.deb' saved
[2446/2446]

Selecting previously unselected package cuda-11-0.
(Reading database ... 124565 files and directories currently installed.)
Preparing to unpack cuda-11-0_11.0.3-1_amd64.deb ...
Unpacking cuda-11-0 (11.0.3-1) ...

dpkg: dependency problems prevent configuration of cuda-11-0:
cuda-11-0 depends on cuda-runtime-11-0 (>= 11.0.3); however:
Package cuda-runtime-11-0 is not installed.
cuda-11-0 depends on cuda-toolkit-11-0 (>= 11.0.3); however:
Package cuda-toolkit-11-0 is not installed.
cuda-11-0 depends on cuda-demo-suite-11-0 (>= 11.0.167); however:
Package cuda-demo-suite-11-0 is not installed.

dpkg: error processing package cuda-11-0 (--install):

dependency problems - leaving unconfigured

Errors were encountered while processing:

cuda-11-0

Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instead
(see apt-key(8)).

Executing: /tmp/apt-key-gpghome.CbYzmmgfRt/gpg.1.sh --fetch-keys https://develop
er.download.nvidia.com/compute/cuda/repos/ubuntu1804/x86_64/7fa2af80.pub

gpg: requesting key from 'https://developer.download.nvidia.com/compute/cuda/repo
s/ubuntu1804/x86_64/7fa2af80.pub'

gpg: key F60F4B3D7FA2AF80: public key "cudatools <cudatools@nvidia.com>"
imported

gpg: Total number processed: 1

gpg: imported: 1

Get:1 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64
InRelease [1,581 B]

Get:2 https://cloud.r-project.org/bin/linux/ubuntu jammy-cran40/ InRelease
[3,626 B]

Get:3 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64
Packages [1,234 kB]

Get:4 https://cloud.r-project.org/bin/linux/ubuntu jammy-cran40/ Packages [62.5
kB]

Get:5 http://security.ubuntu.com/ubuntu jammy-security InRelease [129 kB]

Hit:6 http://archive.ubuntu.com/ubuntu jammy InRelease

Get:7 http://archive.ubuntu.com/ubuntu jammy-updates InRelease [128 kB]

Hit:8 https://ppa.launchpadcontent.net/deadsnakes/ppa/ubuntu jammy InRelease

Get:9 https://r2u.stat.illinois.edu/ubuntu jammy InRelease [6,555 B]

Get:10 http://security.ubuntu.com/ubuntu jammy-security/main amd64 Packages
[2,560 kB]

Hit:11 https://ppa.launchpadcontent.net/graphics-drivers/ppa/ubuntu jammy

```

InRelease
Hit:12 https://ppa.launchpadcontent.net/ubuntugis/ppa/ubuntu jammy InRelease
Get:13 https://r2u.stat.illinois.edu/ubuntu jammy/main amd64 Packages [2,643 kB]
Get:14 http://archive.ubuntu.com/ubuntu jammy-backports InRelease [127 kB]
Get:15 http://security.ubuntu.com/ubuntu jammy-security/universe amd64 Packages
[1,228 kB]
Get:16 http://archive.ubuntu.com/ubuntu jammy-updates/universe amd64 Packages
[1,519 kB]
Get:17 http://archive.ubuntu.com/ubuntu jammy-updates/main amd64 Packages [2,859
kB]
Get:18 https://r2u.stat.illinois.edu/ubuntu jammy/main all Packages [8,607 kB]
Fetched 21.1 MB in 4s (4,951 kB/s)
Reading package lists... Done
W: Skipping acquire of configured file 'main/source/Sources' as repository
'https://r2u.stat.illinois.edu/ubuntu jammy InRelease' does not seem to provide
it (sources.list entry misspelt?)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
You might want to run 'apt --fix-broken install' to correct these.
The following packages have unmet dependencies:
  cuda : Depends: cuda-12-6 (>= 12.6.3) but it is not going to be installed
         Depends: nvidia-open (>= 560.35.05) but it is not going to be installed
  cuda-11-0 : Depends: cuda-runtime-11-0 (>= 11.0.3) but it is not installable
              Depends: cuda-toolkit-11-0 (>= 11.0.3) but it is not installable
              Depends: cuda-demo-suite-11-0 (>= 11.0.167) but it is not
installable
E: Unmet dependencies. Try 'apt --fix-broken install' with no packages (or
specify a solution).
--2025-01-20 03:52:37-- https://developer.download.nvidia.com/compute/redist/cu
dnn/v8.0.5/cudnn-11.0-linux-x64-v8.0.5.39.tgz
Resolving developer.download.nvidia.com (developer.download.nvidia.com)...
23.215.7.4, 23.215.7.28
Connecting to developer.download.nvidia.com
(developer.download.nvidia.com)|23.215.7.4|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1085996495 (1.0G) [application/x-compressed]
Saving to: 'cudnn-11.0-linux-x64-v8.0.5.39.tgz'

cudnn-11.0-linux-x64 100%[=====] 1.01G 3.38MB/s in 5m 9s

2025-01-20 03:57:46 (3.36 MB/s) - 'cudnn-11.0-linux-x64-v8.0.5.39.tgz' saved
[1085996495/1085996495]

cuda/include/cudnn.h
cuda/include/cudnn_adv_infer.h
cuda/include/cudnn_adv_train.h
cuda/include/cudnn_backend.h

```

```

cuda/include/cudnn_cnn_infer.h
cuda/include/cudnn_cnn_train.h
cuda/include/cudnn_ops_infer.h
cuda/include/cudnn_ops_train.h
cuda/include/cudnn_version.h
cuda/NVIDIA_SLA_cuDNN_Support.txt
cuda/lib64/libcudnn.so
cuda/lib64/libcudnn.so.8
cuda/lib64/libcudnn.so.8.0.5
cuda/lib64/libcudnn_adv_infer.so
cuda/lib64/libcudnn_adv_infer.so.8
cuda/lib64/libcudnn_adv_infer.so.8.0.5
cuda/lib64/libcudnn_adv_train.so
cuda/lib64/libcudnn_adv_train.so.8
cuda/lib64/libcudnn_adv_train.so.8.0.5
cuda/lib64/libcudnn_cnn_infer.so
cuda/lib64/libcudnn_cnn_infer.so.8
cuda/lib64/libcudnn_cnn_infer.so.8.0.5
cuda/lib64/libcudnn_cnn_train.so
cuda/lib64/libcudnn_cnn_train.so.8
cuda/lib64/libcudnn_cnn_train.so.8.0.5
cuda/lib64/libcudnn_ops_infer.so
cuda/lib64/libcudnn_ops_infer.so.8
cuda/lib64/libcudnn_ops_infer.so.8.0.5
cuda/lib64/libcudnn_ops_train.so
cuda/lib64/libcudnn_ops_train.so.8
cuda/lib64/libcudnn_ops_train.so.8.0.5
cuda/lib64/libcudnn_static.a
cuda/lib64/libcudnn_static.a

```

Mon Jan 20 03:58:31 2025

```

+-----+
-----+
| NVIDIA-SMI 535.104.05                Driver Version: 535.104.05   CUDA Version:
12.2          |
|-----+-----+-----+
-----+
| GPU  Name                               Persistence-M | Bus-Id        Disp.A | Volatile
Uncorr. ECC |
| Fan  Temp   Perf           Pwr:Usage/Cap |      Memory-Usage | GPU-Util
Compute M.  |
|                |                          |
MIG M.  |
|=====+=====+=====+
=====|
|   0   NVIDIA A100-SXM4-40GB             Off | 00000000:00:04.0 Off |
0 |
| N/A    30C    P0              41W / 400W |      2MiB / 40960MiB |      0%
Default |

```

```

|                                     |
Disabled |
+-----+-----+-----+
-----+

+-----+
-----+
| Processes:
|
| GPU    GI    CI          PID    Type    Process name                      GPU
Memory |
|          ID    ID
Usage   |
|=====
=====|
| No running processes found
|
+-----+
-----+

```

```
[ ]: True
```

```

[2]: from transformers import BertModel, BertTokenizer
import torch
import pandas as pd

# Load the pretrained BERT model and tokenizer (without classification head)
model = BertModel.from_pretrained('bert-base-uncased')
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

def chunk_text_func(input_text, chunk_size=510):
    """
    Split the input text into non-overlapping chunks of specified size.
    The chunking is done by tokenizing the text first, then splitting into
    ↪ chunks.
    """
    tokens = tokenizer.tokenize(input_text)

    chunks = []
    for i in range(0, len(tokens), chunk_size):
        chunk = tokens[i:i + chunk_size]
        chunks.append(chunk)
    return chunks

def extract_cls_embeddings(text, max_length=512):
    """

```

Extracts [CLS] embeddings for each chunk of text.

:param text: The text input to process.
:param chunk_size: The chunk size for splitting long text.
:return: List of CLS embeddings for each chunk of text.
"""

```
chunks = chunk_text_func(text)
cls_embeddings = []
```

```
for chunk in chunks:
```

```
    inputs = tokenizer.encode_plus(
        chunk,
        add_special_tokens=True,
        return_tensors='pt',
        padding='max_length',
        max_length=max_length
    )
```

```
    with torch.no_grad():
        outputs = model(**inputs)
```

```
    cls_embedding = outputs.last_hidden_state[:, 0, :]
    cls_embeddings.append(cls_embedding)
```

```
return cls_embeddings
```

```
def process_texts(df, label):
    """
```

Process texts in a dataframe and return aggregated [CLS] embeddings.

:param df: DataFrame containing the text and label columns.
:param label: The label (0 for beginner, 1 for expert) to filter the
dataframe.
:param chunk_size: The chunk size for splitting long text.
:return: Aggregated mean [CLS] embeddings for the given label.
"""

```
cls_embeddings = []
```

```
for _, row in df[df['label'] == label].iterrows():
    input_text = row['text']
```

```
    if isinstance(input_text, str) and input_text != 'nan':
        cls_embeddings.extend(extract_cls_embeddings(input_text))
    else:
```

```

        # Handle NaN values or invalid text
        print(f"Skipping NaN value in row: {row.name}")

    return torch.mean(torch.stack(cls_embeddings), dim=0)

mean_tensor_beginner = process_texts(train_df, label=0)
mean_tensor_expert = process_texts(train_df, label=1)

'''
# Optionally, print the results or return them as needed
print(f"Mean CLS embedding for beginners: {mean_tensor_beginner}")
print(f"Mean CLS embedding for experts: {mean_tensor_expert}")
'''

```

/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94:
 UserWarning:
 The secret `HF_TOKEN` does not exist in your Colab secrets.
 To authenticate with the Hugging Face Hub, create a token in your settings tab
 (<https://huggingface.co/settings/tokens>), set it as secret in your Google Colab
 and restart your session.
 You will be able to reuse this secret in all of your notebooks.
 Please note that authentication is recommended but still optional to access
 public models or datasets.

```

    warnings.warn(

config.json:   0%|          | 0.00/570 [00:00<?, ?B/s]
model.safetensors:  0%|          | 0.00/440M [00:00<?, ?B/s]
tokenizer_config.json: 0%|          | 0.00/48.0 [00:00<?, ?B/s]
vocab.txt:   0%|          | 0.00/232k [00:00<?, ?B/s]
tokenizer.json:  0%|          | 0.00/466k [00:00<?, ?B/s]

Skipping NaN value in row: 376
Skipping NaN value in row: 470
Skipping NaN value in row: 488
Skipping NaN value in row: 539
Skipping NaN value in row: 545

```

```

[2]: '\n# Optionally, print the results or return them as needed\nprint(f"Mean CLS
embedding for beginners: {mean_tensor_beginner}")\nprint(f"Mean CLS embedding
for experts: {mean_tensor_expert}")\n'

```

```

[3]: from transformers import BertModel, BertTokenizer
import torch

```



```

# Process test texts
test_cls_embeddings = []
for _, test_row in test_df.iterrows():
    input_text = test_row['text']

    if isinstance(input_text, str) and input_text != 'nan': # Check for valid
        ↪text

        cls_embeddings = extract_cls_embeddings(input_text)

        # Aggregate the embeddings (e.g., by averaging the embeddings of all
        ↪chunks)
        mean_cls_embedding = torch.mean(torch.stack(cls_embeddings), dim=0)

        # Append the aggregated [CLS] embedding along with the label
        test_cls_embeddings.append((mean_cls_embedding, test_row['label']))

    else:
        # Handle NaN or invalid text
        print(f"Skipping invalid input_text in test_row: {test_row[0]}") #
        ↪Optional: Print a message for debugging

# Optional: Print the results
print("Aggregated CLS embeddings for the test set:", len(test_cls_embeddings))

```

Aggregated CLS embeddings for the test set: 102

```

[4]: from sklearn.metrics.pairwise import cosine_similarity
correct = 0
beginner_correct = 0
expert_correct = 0

for test_cls_embedding in test_cls_embeddings :

    if test_cls_embedding[1] == 0 :
        truth = "beginner"

    else :
        truth = "expert"

    #TODO : rename refactor and check if using the embeddings
    cls_test_normalized = test_cls_embedding[0].cpu().detach().numpy() / np.
    ↪linalg.norm(test_cls_embedding[0].cpu().detach().numpy())

```

```

cls_beginner_normalized = mean_tensor_beginner.cpu().detach().numpy() / np.
↳linalg.norm(mean_tensor_beginner.cpu().detach().numpy())
similarity_beg = cosine_similarity(cls_test_normalized,
↳cls_beginner_normalized)[0][0]

cls_expert_normalized = mean_tensor_expert.cpu().detach().numpy() / np.linalg.
↳norm(mean_tensor_expert.cpu().detach().numpy())
similarity_expert = cosine_similarity(cls_test_normalized,
↳cls_expert_normalized)[0][0]

if similarity_beg >= similarity_expert and truth == "beginner" :
    beginner_correct += 1
    correct += 1
elif similarity_beg < similarity_expert and truth == "expert" :
    expert_correct += 1
    correct += 1

print("expert correct : ", expert_correct/len(test_df[test_df['label'] ==
↳1]), "\n")
print("beginner correct : ", beginner_correct/len(test_df[test_df['label'] ==
↳0]), "\n")
print("correct : ", correct/len(test_cls_embeddings), "\n")

```

expert correct : 0.6170212765957447

beginner correct : 0.8363636363636363

correct : 0.7352941176470589

```

[5]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.metrics.pairwise import cosine_similarity

# Initialize counters for True Positives (TP), False Positives (FP), and False
↳Negatives (FN)
tp_beg = 0 # True positives for beginner
fp_beg = 0 # False positives for beginner
fn_beg = 0 # False negatives for beginner
tp_exp = 0 # True positives for expert
fp_exp = 0 # False positives for expert
fn_exp = 0 # False negatives for expert

```

```

correct = 0
beginner_correct = 0
expert_correct = 0

# Loop through test_cls_embeddings to classify and update confusion matrix
for test_cls_embedding in test_cls_embeddings:
    if test_cls_embedding[1] == 0:
        truth = "beginner"
    else:
        truth = "expert"

    # Normalize the vectors for cosine similarity
    cls_test_normalized = test_cls_embedding[0].cpu().detach().numpy() / np.
    ↪ linalg.norm(test_cls_embedding[0].cpu().detach().numpy())

    # Compare with beginner's mean vector
    cls_normalized_beg = mean_tensor_beginner.cpu().detach().numpy() / np.
    ↪ linalg.norm(mean_tensor_beginner.cpu().detach().numpy())
    similarity_beg = cosine_similarity(cls_test_normalized.reshape(1, -1),
    ↪ cls_normalized_beg.reshape(1, -1))[0][0]

    # Compare with expert's mean vector
    cls_normalized_exp = mean_tensor_expert.cpu().detach().numpy() / np.linalg.
    ↪ norm(mean_tensor_expert.cpu().detach().numpy())
    similarity_expert = cosine_similarity(cls_test_normalized.reshape(1, -1),
    ↪ cls_normalized_exp.reshape(1, -1))[0][0]

    # Predict the label based on similarity
    predicted = "beginner" if similarity_beg >= similarity_expert else "expert"

    # Update counters for confusion matrix based on prediction vs. truth
    if predicted == "beginner":
        if truth == "beginner":
            tp_beg += 1 # True Positive for beginner
        else:
            fp_beg += 1 # False Positive for beginner
            fn_exp += 1 # False Negative for expert
    elif predicted == "expert":
        if truth == "expert":
            tp_exp += 1 # True Positive for expert
        else:
            fp_exp += 1 # False Positive for expert
            fn_beg += 1 # False Negative for beginner

    # Count correct predictions
    if predicted == truth:
        correct += 1

```

```

    if truth == "beginner":
        beginner_correct += 1
    else:
        expert_correct += 1

# Calculate total True Positives, False Positives, and False Negatives
tp_total = tp_beg + tp_exp
fp_total = fp_beg + fp_exp
fn_total = fn_beg + fn_exp

# Calculate precision, recall, and F1 score for overall classification
precision_total = tp_total / (tp_total + fp_total) if (tp_total + fp_total) > 0
    else 0
recall_total = tp_total / (tp_total + fn_total) if (tp_total + fn_total) > 0
    else 0
f1_total = 2 * (precision_total * recall_total) / (precision_total +
    recall_total) if (precision_total + recall_total) > 0 else 0

# Print the overall F1 score
print(f"Overall F1 score: {f1_total:.4f}")

# Create confusion matrix
conf_matrix = np.array([[tp_beg, fp_exp],
                        [fn_beg, tp_exp]])

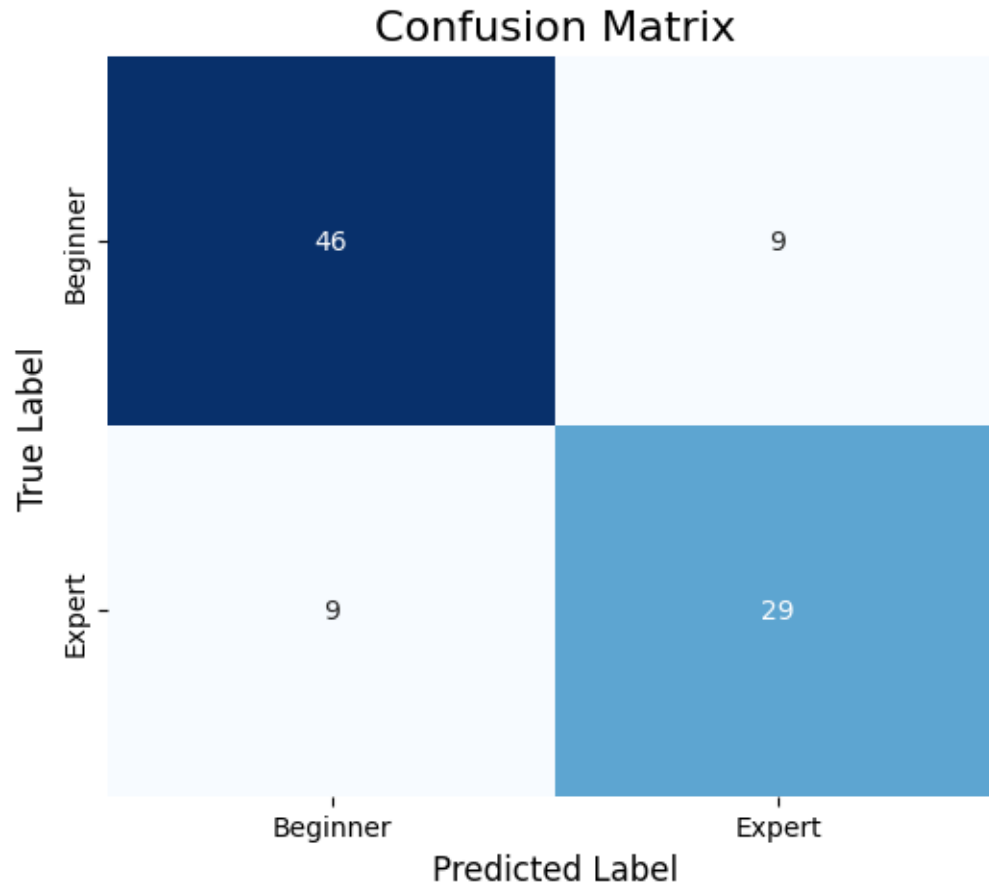
# Plot the confusion matrix using seaborn's heatmap
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
    xticklabels=["Beginner", "Expert"], yticklabels=["Beginner", "Expert"],
    cbar=False)

# Set plot labels and title
plt.title("Confusion Matrix", fontsize=16)
plt.xlabel("Predicted Label", fontsize=12)
plt.ylabel("True Label", fontsize=12)

# Display the plot
plt.show()

```

Overall F1 score: 0.7353



Second method : BERT fine-tuned

```
[6]: import torch
from transformers import BertTokenizer

# Load the tokenizer globally
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

class Dataset(torch.utils.data.Dataset):
    def __init__(self, df, labels, max_length=512):
        """
        Initialize the dataset for BERT processing.

        :param df: pandas DataFrame containing the text and label columns
        :param labels: dictionary or list of labels corresponding to the texts
        :param max_length: maximum sequence length for BERT input
        """
        self.max_length = max_length
        self.labels = []
```

```

self.texts = []
self.document_ids = [] # Store the document ids to link chunks

for idx, text in enumerate(df['text']):
    if isinstance(text, str) and text:
        # Tokenize the text into tokens
        tokens = tokenizer.tokenize(text)

        # Manually chunk the tokenized text
        chunks = self.chunk_tokens(tokens)

        # For each chunk, encode it and append to the dataset
        for chunk in chunks:
            encoded = tokenizer.encode_plus(
                chunk,
                add_special_tokens=True, # Add [CLS] and [SEP]
                ↪automatically
                max_length=max_length,
                padding='max_length', # Pad to max_length
                return_tensors='pt' # Return as PyTorch tensor
            )
            self.texts.append(encoded)
            self.labels.append(labels[df['label'].astype(str).
            ↪iloc[idx]]) # Assuming labels is a dictionary or list
            self.document_ids.append(idx) # Track the document ID

        if len(self.labels) != len(self.texts):
            raise ValueError("Mismatch between number of texts and labels.")

def chunk_tokens(self, tokens, chunk_size=510):
    """Chunk tokenized text into smaller parts."""
    chunks = []
    for i in range(0, len(tokens), chunk_size):
        chunk = tokens[i:i + chunk_size]
        chunks.append(chunk)
    return chunks

def __len__(self):
    return len(self.texts)

def __getitem__(self, idx):
    text = self.texts[idx]
    label = self.labels[idx]
    document_id = self.document_ids[idx]
    return text['input_ids'].squeeze(0), text['attention_mask'].squeeze(0), ↪
    ↪label, document_id

```

```

[7]: import torch
import torch.nn as nn
from transformers import BertForSequenceClassification

class BertClassifier(nn.Module):
    def __init__(self, num_classes, dropout=0.1):
        super(BertClassifier, self).__init__()

        # Use BertForSequenceClassification which includes BERT model +
        ↪ classification head
        self.bert = BertForSequenceClassification.
        ↪ from_pretrained('bert-base-uncased', num_labels=num_classes,
        ↪ output_hidden_states=True)
        self.dropout = nn.Dropout(dropout) # Optional dropout layer
        self.cls_embeddings = None
        self.logits = None

    def forward(self, input_ids, attention_mask):
        """
        Forward pass through the BERT model and aggregate CLS embeddings by
        ↪ document.

        :param input_ids: Tensor of input token IDs (batch_size x seq_len)
        :param attention_mask: Tensor of attention masks (batch_size x seq_len)
        :param document_ids: Tensor of document IDs (batch_size), linking
        ↪ chunks from the same document
        :return: logits: Classification logits
        """
        # Forward pass through BERT model
        outputs = self.bert(input_ids=input_ids, attention_mask=attention_mask,
        ↪ output_hidden_states=True)

        # Extract the last hidden state (the last layer hidden states)
        last_hidden_state = outputs.hidden_states[-1] # Shape: (batch_size,
        ↪ sequence_length, hidden_size)

        # Get the CLS embedding for each chunk (first token in each sequence)
        cls_embeddings = last_hidden_state[:, 0, :] # Shape: (batch_size,
        ↪ hidden_size)

        # Apply dropout
        cls_embeddings = self.dropout(cls_embeddings)

        self.cls_embeddings = cls_embeddings

```

```

        logits = self.bert.classifier(cls_embeddings) # Apply classification
↪head
        self.logits = logits

        return logits

```

```

[8]: import torch
from torch.utils.data import DataLoader
from torch.nn.utils.rnn import pad_sequence
from torch.optim import Adam
from tqdm import tqdm
import torch.nn as nn
from transformers import BertTokenizer

# Load the tokenizer globally
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

class EarlyStopping:
    def __init__(self, patience=5, verbose=False):
        self.patience = patience
        self.verbose = verbose
        self.best_loss = float('inf')
        self.counter = 0
        self.early_stop = False

    def __call__(self, val_loss):
        if val_loss < self.best_loss:
            self.best_loss = val_loss
            self.counter = 0
        else:
            self.counter += 1
            if self.counter >= self.patience:
                self.early_stop = True
                if self.verbose:
                    print(f"Early stopping triggered after {self.patience}
↪epochs without improvement.")

def train(model, train_data, val_data, labels, learning_rate, epochs):
    # Initialize datasets
    train_dataset = Dataset(train_data, labels)
    val_dataset = Dataset(val_data, labels)

    def collate_fn(batch):

        input_ids = pad_sequence([item[0] for item in batch], batch_first=True,
↪padding_value=tokenizer.pad_token_id)

```



```

        attention_mask = pad_sequence([item[1] for item in batch],␣
↪batch_first=True, padding_value=0)
        labels = torch.tensor([item[2] for item in batch])
        document_ids = torch.tensor([item[3] for item in batch])
        return {'input_ids': input_ids, 'attention_mask': attention_mask,␣
↪'document_ids': document_ids}, labels

    # DataLoader
    train_dataloader = DataLoader(train_dataset, batch_size=2, shuffle=True,␣
↪collate_fn=collate_fn, num_workers=2)
    val_dataloader = DataLoader(val_dataset, batch_size=2,␣
↪collate_fn=collate_fn, num_workers=2)

    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    model.to(device)

    criterion = nn.CrossEntropyLoss().to(device)
    optimizer = Adam(model.parameters(), lr=learning_rate)

    cls_embeddings_with_text_train = []
    cls_embeddings_with_text_val = []

    early_stopping = EarlyStopping(patience=5, verbose=True)

    # Gradient accumulation steps
    accumulation_steps = 4 # Adjust this based on your needs

    for epoch_num in range(epochs):
        total_acc_train, total_loss_train = 0, 0
        model.train()

        for i, (train_input, train_label) in enumerate(tqdm(train_dataloader)):
            train_label = train_label.to(device)
            input_ids = train_input['input_ids'].to(device)
            attention_mask = train_input['attention_mask'].to(device)
            document_ids = train_input['document_ids'].to(device)

            # Forward pass
            output = model(input_ids, attention_mask)
            batch_loss = criterion(output, train_label)
            total_loss_train += batch_loss.item()

            # Calculate accuracy
            total_correct_train = (output.argmax(dim=1) == train_label).sum().
↪item()

            total_acc_train += total_correct_train

```

```

        # Normalize loss to account for accumulation
        batch_loss = batch_loss / accumulation_steps
        batch_loss.backward()

    if (i + 1) % accumulation_steps == 0:
        optimizer.step()
        optimizer.zero_grad()

    # Store CLS embeddings and text for further analysis
    # TODO. : double check the below function
    cls_embeddings_with_text_train.extend(zip(
        [tokenizer.decode(ids, skip_special_tokens=True) for ids in ↪
input_ids],
        model.cls_embeddings.detach().cpu().numpy(),
        train_label.cpu().numpy()
    ))

    # Validation loop
    total_acc_val, total_loss_val = 0, 0
    model.eval()

    with torch.no_grad():
        for val_input, val_label in val_dataloader:
            val_label = val_label.to(device)
            input_ids = val_input['input_ids'].to(device)
            attention_mask = val_input['attention_mask'].to(device)
            document_ids = val_input['document_ids'].to(device)

            output = model(input_ids, attention_mask)
            batch_loss = criterion(output, val_label)
            total_loss_val += batch_loss.item()

            # Calculate accuracy for validation
            total_correct_val = (output.argmax(dim=1) == val_label).sum().
↪item()

            total_acc_val += total_correct_val

        # Store CLS embeddings and text for validation
        cls_embeddings_with_text_val.extend(zip(
            [tokenizer.decode(ids, skip_special_tokens=True) for ids in ↪
input_ids],
            model.cls_embeddings.detach().cpu().numpy(),
            val_label.cpu().numpy()
        ))

    # Calculate average losses and accuracies
    avg_train_loss = total_loss_train / len(train_dataloader)

```

```

        avg_train_acc = total_acc_train / (len(train_dataloader) *
        ↪train_dataloader.batch_size)
        avg_val_loss = total_loss_val / len(val_dataloader)
        avg_val_acc = total_acc_val / (len(val_dataloader) * val_dataloader.
        ↪batch_size)

        print(f'Epochs: {epoch_num + 1} | Train Loss: {avg_train_loss:.3f} |
        ↪Train Accuracy: {avg_train_acc:.3f} | Val Loss: {avg_val_loss:.3f} | Val
        ↪Accuracy: {avg_val_acc:.3f}')

        # Print which parameters were updated
        print("Updated parameters during this epoch:")
        updated_params = []
        for name, param in model.named_parameters():
            if param.grad is not None: # Only parameters that were updated
        ↪will have gradients
                updated_params.append(name)
        print(f"Parameters updated: {updated_params}")

        # Check for early stopping
        early_stopping(avg_val_loss)
        if early_stopping.early_stop:
            print("Stopping training early!")
            break

    return cls_embeddings_with_text_train, cls_embeddings_with_text_val

```

```

[9]: EPOCHS = 20
    num_classes = 2
    model = BertClassifier(num_classes=num_classes)
    LR = 1e-05

    labels = {
        '0': 0,
        '1': 1
    }

    cls_embeddings_with_text_train,cls_embeddings_with_text_val = train(model,
    ↪train_df, val_df, labels, LR, EPOCHS)

```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized:

```
['classifier.bias', 'classifier.weight']
```

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
100% | 1054/1054 [01:06<00:00, 15.88it/s]
```

Epochs: 1 | Train Loss: 0.289 | Train Accuracy: 0.880 | Val Loss: 0.172 | Val Accuracy: 0.934

Updated parameters during this epoch:

Parameters updated: ['bert.bert.embeddings.word_embeddings.weight',
'bert.bert.embeddings.position_embeddings.weight',
'bert.bert.embeddings.token_type_embeddings.weight',
'bert.bert.embeddings.LayerNorm.weight', 'bert.bert.embeddings.LayerNorm.bias',
'bert.bert.encoder.layer.0.attention.self.query.weight',
'bert.bert.encoder.layer.0.attention.self.query.bias',
'bert.bert.encoder.layer.0.attention.self.key.weight',
'bert.bert.encoder.layer.0.attention.self.key.bias',
'bert.bert.encoder.layer.0.attention.self.value.weight',
'bert.bert.encoder.layer.0.attention.self.value.bias',
'bert.bert.encoder.layer.0.attention.output.dense.weight',
'bert.bert.encoder.layer.0.attention.output.dense.bias',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.0.intermediate.dense.weight',
'bert.bert.encoder.layer.0.intermediate.dense.bias',
'bert.bert.encoder.layer.0.output.dense.weight',
'bert.bert.encoder.layer.0.output.dense.bias',
'bert.bert.encoder.layer.0.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.attention.self.query.weight',
'bert.bert.encoder.layer.1.attention.self.query.bias',
'bert.bert.encoder.layer.1.attention.self.key.weight',
'bert.bert.encoder.layer.1.attention.self.key.bias',
'bert.bert.encoder.layer.1.attention.self.value.weight',
'bert.bert.encoder.layer.1.attention.self.value.bias',
'bert.bert.encoder.layer.1.attention.output.dense.weight',
'bert.bert.encoder.layer.1.attention.output.dense.bias',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.intermediate.dense.weight',
'bert.bert.encoder.layer.1.intermediate.dense.bias',
'bert.bert.encoder.layer.1.output.dense.weight',
'bert.bert.encoder.layer.1.output.dense.bias',
'bert.bert.encoder.layer.1.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.attention.self.query.weight',
'bert.bert.encoder.layer.2.attention.self.query.bias',
'bert.bert.encoder.layer.2.attention.self.key.weight',
'bert.bert.encoder.layer.2.attention.self.key.bias',
'bert.bert.encoder.layer.2.attention.self.value.weight',
'bert.bert.encoder.layer.2.attention.self.value.bias',
'bert.bert.encoder.layer.2.attention.output.dense.weight',
'bert.bert.encoder.layer.2.attention.output.dense.bias',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.weight',

'bert.bert.encoder.layer.2.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.intermediate.dense.weight',
'bert.bert.encoder.layer.2.intermediate.dense.bias',
'bert.bert.encoder.layer.2.output.dense.weight',
'bert.bert.encoder.layer.2.output.dense.bias',
'bert.bert.encoder.layer.2.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.attention.self.query.weight',
'bert.bert.encoder.layer.3.attention.self.query.bias',
'bert.bert.encoder.layer.3.attention.self.key.weight',
'bert.bert.encoder.layer.3.attention.self.key.bias',
'bert.bert.encoder.layer.3.attention.self.value.weight',
'bert.bert.encoder.layer.3.attention.self.value.bias',
'bert.bert.encoder.layer.3.attention.output.dense.weight',
'bert.bert.encoder.layer.3.attention.output.dense.bias',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.intermediate.dense.weight',
'bert.bert.encoder.layer.3.intermediate.dense.bias',
'bert.bert.encoder.layer.3.output.dense.weight',
'bert.bert.encoder.layer.3.output.dense.bias',
'bert.bert.encoder.layer.3.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.attention.self.query.weight',
'bert.bert.encoder.layer.4.attention.self.query.bias',
'bert.bert.encoder.layer.4.attention.self.key.weight',
'bert.bert.encoder.layer.4.attention.self.key.bias',
'bert.bert.encoder.layer.4.attention.self.value.weight',
'bert.bert.encoder.layer.4.attention.self.value.bias',
'bert.bert.encoder.layer.4.attention.output.dense.weight',
'bert.bert.encoder.layer.4.attention.output.dense.bias',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.intermediate.dense.weight',
'bert.bert.encoder.layer.4.intermediate.dense.bias',
'bert.bert.encoder.layer.4.output.dense.weight',
'bert.bert.encoder.layer.4.output.dense.bias',
'bert.bert.encoder.layer.4.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.attention.self.query.weight',
'bert.bert.encoder.layer.5.attention.self.query.bias',
'bert.bert.encoder.layer.5.attention.self.key.weight',
'bert.bert.encoder.layer.5.attention.self.key.bias',
'bert.bert.encoder.layer.5.attention.self.value.weight',
'bert.bert.encoder.layer.5.attention.self.value.bias',
'bert.bert.encoder.layer.5.attention.output.dense.weight',
'bert.bert.encoder.layer.5.attention.output.dense.bias',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.weight',

'bert.bert.encoder.layer.5.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.intermediate.dense.weight',
'bert.bert.encoder.layer.5.intermediate.dense.bias',
'bert.bert.encoder.layer.5.output.dense.weight',
'bert.bert.encoder.layer.5.output.dense.bias',
'bert.bert.encoder.layer.5.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.attention.self.query.weight',
'bert.bert.encoder.layer.6.attention.self.query.bias',
'bert.bert.encoder.layer.6.attention.self.key.weight',
'bert.bert.encoder.layer.6.attention.self.key.bias',
'bert.bert.encoder.layer.6.attention.self.value.weight',
'bert.bert.encoder.layer.6.attention.self.value.bias',
'bert.bert.encoder.layer.6.attention.output.dense.weight',
'bert.bert.encoder.layer.6.attention.output.dense.bias',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.intermediate.dense.weight',
'bert.bert.encoder.layer.6.intermediate.dense.bias',
'bert.bert.encoder.layer.6.output.dense.weight',
'bert.bert.encoder.layer.6.output.dense.bias',
'bert.bert.encoder.layer.6.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.attention.self.query.weight',
'bert.bert.encoder.layer.7.attention.self.query.bias',
'bert.bert.encoder.layer.7.attention.self.key.weight',
'bert.bert.encoder.layer.7.attention.self.key.bias',
'bert.bert.encoder.layer.7.attention.self.value.weight',
'bert.bert.encoder.layer.7.attention.self.value.bias',
'bert.bert.encoder.layer.7.attention.output.dense.weight',
'bert.bert.encoder.layer.7.attention.output.dense.bias',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.intermediate.dense.weight',
'bert.bert.encoder.layer.7.intermediate.dense.bias',
'bert.bert.encoder.layer.7.output.dense.weight',
'bert.bert.encoder.layer.7.output.dense.bias',
'bert.bert.encoder.layer.7.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.attention.self.query.weight',
'bert.bert.encoder.layer.8.attention.self.query.bias',
'bert.bert.encoder.layer.8.attention.self.key.weight',
'bert.bert.encoder.layer.8.attention.self.key.bias',
'bert.bert.encoder.layer.8.attention.self.value.weight',
'bert.bert.encoder.layer.8.attention.self.value.bias',
'bert.bert.encoder.layer.8.attention.output.dense.weight',
'bert.bert.encoder.layer.8.attention.output.dense.bias',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.weight',

'bert.bert.encoder.layer.8.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.intermediate.dense.weight',
'bert.bert.encoder.layer.8.intermediate.dense.bias',
'bert.bert.encoder.layer.8.output.dense.weight',
'bert.bert.encoder.layer.8.output.dense.bias',
'bert.bert.encoder.layer.8.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.attention.self.query.weight',
'bert.bert.encoder.layer.9.attention.self.query.bias',
'bert.bert.encoder.layer.9.attention.self.key.weight',
'bert.bert.encoder.layer.9.attention.self.key.bias',
'bert.bert.encoder.layer.9.attention.self.value.weight',
'bert.bert.encoder.layer.9.attention.self.value.bias',
'bert.bert.encoder.layer.9.attention.output.dense.weight',
'bert.bert.encoder.layer.9.attention.output.dense.bias',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.intermediate.dense.weight',
'bert.bert.encoder.layer.9.intermediate.dense.bias',
'bert.bert.encoder.layer.9.output.dense.weight',
'bert.bert.encoder.layer.9.output.dense.bias',
'bert.bert.encoder.layer.9.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.attention.self.query.weight',
'bert.bert.encoder.layer.10.attention.self.query.bias',
'bert.bert.encoder.layer.10.attention.self.key.weight',
'bert.bert.encoder.layer.10.attention.self.key.bias',
'bert.bert.encoder.layer.10.attention.self.value.weight',
'bert.bert.encoder.layer.10.attention.self.value.bias',
'bert.bert.encoder.layer.10.attention.output.dense.weight',
'bert.bert.encoder.layer.10.attention.output.dense.bias',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.intermediate.dense.weight',
'bert.bert.encoder.layer.10.intermediate.dense.bias',
'bert.bert.encoder.layer.10.output.dense.weight',
'bert.bert.encoder.layer.10.output.dense.bias',
'bert.bert.encoder.layer.10.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.attention.self.query.weight',
'bert.bert.encoder.layer.11.attention.self.query.bias',
'bert.bert.encoder.layer.11.attention.self.key.weight',
'bert.bert.encoder.layer.11.attention.self.key.bias',
'bert.bert.encoder.layer.11.attention.self.value.weight',
'bert.bert.encoder.layer.11.attention.self.value.bias',
'bert.bert.encoder.layer.11.attention.output.dense.weight',
'bert.bert.encoder.layer.11.attention.output.dense.bias',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.weight',

```
'bert.bert.encoder.layer.11.attention.output.LayerNorm.bias',  
'bert.bert.encoder.layer.11.intermediate.dense.weight',  
'bert.bert.encoder.layer.11.intermediate.dense.bias',  
'bert.bert.encoder.layer.11.output.dense.weight',  
'bert.bert.encoder.layer.11.output.dense.bias',  
'bert.bert.encoder.layer.11.output.LayerNorm.weight',  
'bert.bert.encoder.layer.11.output.LayerNorm.bias', 'bert.classifier.weight',  
'bert.classifier.bias']
```

100%| | 1054/1054 [01:05<00:00, 16.19it/s]

Epochs: 2 | Train Loss: 0.120 | Train Accuracy: 0.957 | Val Loss: 0.169 | Val Accuracy: 0.931

Updated parameters during this epoch:

Parameters updated: ['bert.bert.embeddings.word_embeddings.weight',
'bert.bert.embeddings.position_embeddings.weight',
'bert.bert.embeddings.token_type_embeddings.weight',
'bert.bert.embeddings.LayerNorm.weight', 'bert.bert.embeddings.LayerNorm.bias',
'bert.bert.encoder.layer.0.attention.self.query.weight',
'bert.bert.encoder.layer.0.attention.self.query.bias',
'bert.bert.encoder.layer.0.attention.self.key.weight',
'bert.bert.encoder.layer.0.attention.self.key.bias',
'bert.bert.encoder.layer.0.attention.self.value.weight',
'bert.bert.encoder.layer.0.attention.self.value.bias',
'bert.bert.encoder.layer.0.attention.output.dense.weight',
'bert.bert.encoder.layer.0.attention.output.dense.bias',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.0.intermediate.dense.weight',
'bert.bert.encoder.layer.0.intermediate.dense.bias',
'bert.bert.encoder.layer.0.output.dense.weight',
'bert.bert.encoder.layer.0.output.dense.bias',
'bert.bert.encoder.layer.0.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.attention.self.query.weight',
'bert.bert.encoder.layer.1.attention.self.query.bias',
'bert.bert.encoder.layer.1.attention.self.key.weight',
'bert.bert.encoder.layer.1.attention.self.key.bias',
'bert.bert.encoder.layer.1.attention.self.value.weight',
'bert.bert.encoder.layer.1.attention.self.value.bias',
'bert.bert.encoder.layer.1.attention.output.dense.weight',
'bert.bert.encoder.layer.1.attention.output.dense.bias',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.intermediate.dense.weight',
'bert.bert.encoder.layer.1.intermediate.dense.bias',
'bert.bert.encoder.layer.1.output.dense.weight',
'bert.bert.encoder.layer.1.output.dense.bias',
'bert.bert.encoder.layer.1.output.LayerNorm.weight',

'bert.bert.encoder.layer.1.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.attention.self.query.weight',
'bert.bert.encoder.layer.2.attention.self.query.bias',
'bert.bert.encoder.layer.2.attention.self.key.weight',
'bert.bert.encoder.layer.2.attention.self.key.bias',
'bert.bert.encoder.layer.2.attention.self.value.weight',
'bert.bert.encoder.layer.2.attention.self.value.bias',
'bert.bert.encoder.layer.2.attention.output.dense.weight',
'bert.bert.encoder.layer.2.attention.output.dense.bias',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.intermediate.dense.weight',
'bert.bert.encoder.layer.2.intermediate.dense.bias',
'bert.bert.encoder.layer.2.output.dense.weight',
'bert.bert.encoder.layer.2.output.dense.bias',
'bert.bert.encoder.layer.2.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.attention.self.query.weight',
'bert.bert.encoder.layer.3.attention.self.query.bias',
'bert.bert.encoder.layer.3.attention.self.key.weight',
'bert.bert.encoder.layer.3.attention.self.key.bias',
'bert.bert.encoder.layer.3.attention.self.value.weight',
'bert.bert.encoder.layer.3.attention.self.value.bias',
'bert.bert.encoder.layer.3.attention.output.dense.weight',
'bert.bert.encoder.layer.3.attention.output.dense.bias',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.intermediate.dense.weight',
'bert.bert.encoder.layer.3.intermediate.dense.bias',
'bert.bert.encoder.layer.3.output.dense.weight',
'bert.bert.encoder.layer.3.output.dense.bias',
'bert.bert.encoder.layer.3.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.attention.self.query.weight',
'bert.bert.encoder.layer.4.attention.self.query.bias',
'bert.bert.encoder.layer.4.attention.self.key.weight',
'bert.bert.encoder.layer.4.attention.self.key.bias',
'bert.bert.encoder.layer.4.attention.self.value.weight',
'bert.bert.encoder.layer.4.attention.self.value.bias',
'bert.bert.encoder.layer.4.attention.output.dense.weight',
'bert.bert.encoder.layer.4.attention.output.dense.bias',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.intermediate.dense.weight',
'bert.bert.encoder.layer.4.intermediate.dense.bias',
'bert.bert.encoder.layer.4.output.dense.weight',
'bert.bert.encoder.layer.4.output.dense.bias',
'bert.bert.encoder.layer.4.output.LayerNorm.weight',

'bert.bert.encoder.layer.4.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.attention.self.query.weight',
'bert.bert.encoder.layer.5.attention.self.query.bias',
'bert.bert.encoder.layer.5.attention.self.key.weight',
'bert.bert.encoder.layer.5.attention.self.key.bias',
'bert.bert.encoder.layer.5.attention.self.value.weight',
'bert.bert.encoder.layer.5.attention.self.value.bias',
'bert.bert.encoder.layer.5.attention.output.dense.weight',
'bert.bert.encoder.layer.5.attention.output.dense.bias',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.intermediate.dense.weight',
'bert.bert.encoder.layer.5.intermediate.dense.bias',
'bert.bert.encoder.layer.5.output.dense.weight',
'bert.bert.encoder.layer.5.output.dense.bias',
'bert.bert.encoder.layer.5.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.attention.self.query.weight',
'bert.bert.encoder.layer.6.attention.self.query.bias',
'bert.bert.encoder.layer.6.attention.self.key.weight',
'bert.bert.encoder.layer.6.attention.self.key.bias',
'bert.bert.encoder.layer.6.attention.self.value.weight',
'bert.bert.encoder.layer.6.attention.self.value.bias',
'bert.bert.encoder.layer.6.attention.output.dense.weight',
'bert.bert.encoder.layer.6.attention.output.dense.bias',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.intermediate.dense.weight',
'bert.bert.encoder.layer.6.intermediate.dense.bias',
'bert.bert.encoder.layer.6.output.dense.weight',
'bert.bert.encoder.layer.6.output.dense.bias',
'bert.bert.encoder.layer.6.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.attention.self.query.weight',
'bert.bert.encoder.layer.7.attention.self.query.bias',
'bert.bert.encoder.layer.7.attention.self.key.weight',
'bert.bert.encoder.layer.7.attention.self.key.bias',
'bert.bert.encoder.layer.7.attention.self.value.weight',
'bert.bert.encoder.layer.7.attention.self.value.bias',
'bert.bert.encoder.layer.7.attention.output.dense.weight',
'bert.bert.encoder.layer.7.attention.output.dense.bias',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.intermediate.dense.weight',
'bert.bert.encoder.layer.7.intermediate.dense.bias',
'bert.bert.encoder.layer.7.output.dense.weight',
'bert.bert.encoder.layer.7.output.dense.bias',
'bert.bert.encoder.layer.7.output.LayerNorm.weight',

'bert.bert.encoder.layer.7.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.attention.self.query.weight',
'bert.bert.encoder.layer.8.attention.self.query.bias',
'bert.bert.encoder.layer.8.attention.self.key.weight',
'bert.bert.encoder.layer.8.attention.self.key.bias',
'bert.bert.encoder.layer.8.attention.self.value.weight',
'bert.bert.encoder.layer.8.attention.self.value.bias',
'bert.bert.encoder.layer.8.attention.output.dense.weight',
'bert.bert.encoder.layer.8.attention.output.dense.bias',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.intermediate.dense.weight',
'bert.bert.encoder.layer.8.intermediate.dense.bias',
'bert.bert.encoder.layer.8.output.dense.weight',
'bert.bert.encoder.layer.8.output.dense.bias',
'bert.bert.encoder.layer.8.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.attention.self.query.weight',
'bert.bert.encoder.layer.9.attention.self.query.bias',
'bert.bert.encoder.layer.9.attention.self.key.weight',
'bert.bert.encoder.layer.9.attention.self.key.bias',
'bert.bert.encoder.layer.9.attention.self.value.weight',
'bert.bert.encoder.layer.9.attention.self.value.bias',
'bert.bert.encoder.layer.9.attention.output.dense.weight',
'bert.bert.encoder.layer.9.attention.output.dense.bias',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.intermediate.dense.weight',
'bert.bert.encoder.layer.9.intermediate.dense.bias',
'bert.bert.encoder.layer.9.output.dense.weight',
'bert.bert.encoder.layer.9.output.dense.bias',
'bert.bert.encoder.layer.9.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.attention.self.query.weight',
'bert.bert.encoder.layer.10.attention.self.query.bias',
'bert.bert.encoder.layer.10.attention.self.key.weight',
'bert.bert.encoder.layer.10.attention.self.key.bias',
'bert.bert.encoder.layer.10.attention.self.value.weight',
'bert.bert.encoder.layer.10.attention.self.value.bias',
'bert.bert.encoder.layer.10.attention.output.dense.weight',
'bert.bert.encoder.layer.10.attention.output.dense.bias',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.intermediate.dense.weight',
'bert.bert.encoder.layer.10.intermediate.dense.bias',
'bert.bert.encoder.layer.10.output.dense.weight',
'bert.bert.encoder.layer.10.output.dense.bias',
'bert.bert.encoder.layer.10.output.LayerNorm.weight',

```
'bert.bert.encoder.layer.10.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.attention.self.query.weight',
'bert.bert.encoder.layer.11.attention.self.query.bias',
'bert.bert.encoder.layer.11.attention.self.key.weight',
'bert.bert.encoder.layer.11.attention.self.key.bias',
'bert.bert.encoder.layer.11.attention.self.value.weight',
'bert.bert.encoder.layer.11.attention.self.value.bias',
'bert.bert.encoder.layer.11.attention.output.dense.weight',
'bert.bert.encoder.layer.11.attention.output.dense.bias',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.intermediate.dense.weight',
'bert.bert.encoder.layer.11.intermediate.dense.bias',
'bert.bert.encoder.layer.11.output.dense.weight',
'bert.bert.encoder.layer.11.output.dense.bias',
'bert.bert.encoder.layer.11.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.output.LayerNorm.bias', 'bert.classifier.weight',
'bert.classifier.bias']
```

100%| | 1054/1054 [01:05<00:00, 16.20it/s]

Epochs: 3 | Train Loss: 0.048 | Train Accuracy: 0.986 | Val Loss: 0.182 | Val Accuracy: 0.949

Updated parameters during this epoch:

Parameters updated: ['bert.bert.embeddings.word_embeddings.weight',
'bert.bert.embeddings.position_embeddings.weight',
'bert.bert.embeddings.token_type_embeddings.weight',
'bert.bert.embeddings.LayerNorm.weight', 'bert.bert.embeddings.LayerNorm.bias',
'bert.bert.encoder.layer.0.attention.self.query.weight',
'bert.bert.encoder.layer.0.attention.self.query.bias',
'bert.bert.encoder.layer.0.attention.self.key.weight',
'bert.bert.encoder.layer.0.attention.self.key.bias',
'bert.bert.encoder.layer.0.attention.self.value.weight',
'bert.bert.encoder.layer.0.attention.self.value.bias',
'bert.bert.encoder.layer.0.attention.output.dense.weight',
'bert.bert.encoder.layer.0.attention.output.dense.bias',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.0.intermediate.dense.weight',
'bert.bert.encoder.layer.0.intermediate.dense.bias',
'bert.bert.encoder.layer.0.output.dense.weight',
'bert.bert.encoder.layer.0.output.dense.bias',
'bert.bert.encoder.layer.0.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.attention.self.query.weight',
'bert.bert.encoder.layer.1.attention.self.query.bias',
'bert.bert.encoder.layer.1.attention.self.key.weight',
'bert.bert.encoder.layer.1.attention.self.key.bias',
'bert.bert.encoder.layer.1.attention.self.value.weight',

'bert.bert.encoder.layer.1.attention.self.value.bias',
'bert.bert.encoder.layer.1.attention.output.dense.weight',
'bert.bert.encoder.layer.1.attention.output.dense.bias',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.intermediate.dense.weight',
'bert.bert.encoder.layer.1.intermediate.dense.bias',
'bert.bert.encoder.layer.1.output.dense.weight',
'bert.bert.encoder.layer.1.output.dense.bias',
'bert.bert.encoder.layer.1.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.attention.self.query.weight',
'bert.bert.encoder.layer.2.attention.self.query.bias',
'bert.bert.encoder.layer.2.attention.self.key.weight',
'bert.bert.encoder.layer.2.attention.self.key.bias',
'bert.bert.encoder.layer.2.attention.self.value.weight',
'bert.bert.encoder.layer.2.attention.self.value.bias',
'bert.bert.encoder.layer.2.attention.output.dense.weight',
'bert.bert.encoder.layer.2.attention.output.dense.bias',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.intermediate.dense.weight',
'bert.bert.encoder.layer.2.intermediate.dense.bias',
'bert.bert.encoder.layer.2.output.dense.weight',
'bert.bert.encoder.layer.2.output.dense.bias',
'bert.bert.encoder.layer.2.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.attention.self.query.weight',
'bert.bert.encoder.layer.3.attention.self.query.bias',
'bert.bert.encoder.layer.3.attention.self.key.weight',
'bert.bert.encoder.layer.3.attention.self.key.bias',
'bert.bert.encoder.layer.3.attention.self.value.weight',
'bert.bert.encoder.layer.3.attention.self.value.bias',
'bert.bert.encoder.layer.3.attention.output.dense.weight',
'bert.bert.encoder.layer.3.attention.output.dense.bias',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.intermediate.dense.weight',
'bert.bert.encoder.layer.3.intermediate.dense.bias',
'bert.bert.encoder.layer.3.output.dense.weight',
'bert.bert.encoder.layer.3.output.dense.bias',
'bert.bert.encoder.layer.3.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.attention.self.query.weight',
'bert.bert.encoder.layer.4.attention.self.query.bias',
'bert.bert.encoder.layer.4.attention.self.key.weight',
'bert.bert.encoder.layer.4.attention.self.key.bias',
'bert.bert.encoder.layer.4.attention.self.value.weight',

'bert.bert.encoder.layer.4.attention.self.value.bias',
'bert.bert.encoder.layer.4.attention.output.dense.weight',
'bert.bert.encoder.layer.4.attention.output.dense.bias',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.intermediate.dense.weight',
'bert.bert.encoder.layer.4.intermediate.dense.bias',
'bert.bert.encoder.layer.4.output.dense.weight',
'bert.bert.encoder.layer.4.output.dense.bias',
'bert.bert.encoder.layer.4.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.attention.self.query.weight',
'bert.bert.encoder.layer.5.attention.self.query.bias',
'bert.bert.encoder.layer.5.attention.self.key.weight',
'bert.bert.encoder.layer.5.attention.self.key.bias',
'bert.bert.encoder.layer.5.attention.self.value.weight',
'bert.bert.encoder.layer.5.attention.self.value.bias',
'bert.bert.encoder.layer.5.attention.output.dense.weight',
'bert.bert.encoder.layer.5.attention.output.dense.bias',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.intermediate.dense.weight',
'bert.bert.encoder.layer.5.intermediate.dense.bias',
'bert.bert.encoder.layer.5.output.dense.weight',
'bert.bert.encoder.layer.5.output.dense.bias',
'bert.bert.encoder.layer.5.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.attention.self.query.weight',
'bert.bert.encoder.layer.6.attention.self.query.bias',
'bert.bert.encoder.layer.6.attention.self.key.weight',
'bert.bert.encoder.layer.6.attention.self.key.bias',
'bert.bert.encoder.layer.6.attention.self.value.weight',
'bert.bert.encoder.layer.6.attention.self.value.bias',
'bert.bert.encoder.layer.6.attention.output.dense.weight',
'bert.bert.encoder.layer.6.attention.output.dense.bias',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.intermediate.dense.weight',
'bert.bert.encoder.layer.6.intermediate.dense.bias',
'bert.bert.encoder.layer.6.output.dense.weight',
'bert.bert.encoder.layer.6.output.dense.bias',
'bert.bert.encoder.layer.6.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.attention.self.query.weight',
'bert.bert.encoder.layer.7.attention.self.query.bias',
'bert.bert.encoder.layer.7.attention.self.key.weight',
'bert.bert.encoder.layer.7.attention.self.key.bias',
'bert.bert.encoder.layer.7.attention.self.value.weight',

'bert.bert.encoder.layer.7.attention.self.value.bias',
'bert.bert.encoder.layer.7.attention.output.dense.weight',
'bert.bert.encoder.layer.7.attention.output.dense.bias',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.intermediate.dense.weight',
'bert.bert.encoder.layer.7.intermediate.dense.bias',
'bert.bert.encoder.layer.7.output.dense.weight',
'bert.bert.encoder.layer.7.output.dense.bias',
'bert.bert.encoder.layer.7.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.attention.self.query.weight',
'bert.bert.encoder.layer.8.attention.self.query.bias',
'bert.bert.encoder.layer.8.attention.self.key.weight',
'bert.bert.encoder.layer.8.attention.self.key.bias',
'bert.bert.encoder.layer.8.attention.self.value.weight',
'bert.bert.encoder.layer.8.attention.self.value.bias',
'bert.bert.encoder.layer.8.attention.output.dense.weight',
'bert.bert.encoder.layer.8.attention.output.dense.bias',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.intermediate.dense.weight',
'bert.bert.encoder.layer.8.intermediate.dense.bias',
'bert.bert.encoder.layer.8.output.dense.weight',
'bert.bert.encoder.layer.8.output.dense.bias',
'bert.bert.encoder.layer.8.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.attention.self.query.weight',
'bert.bert.encoder.layer.9.attention.self.query.bias',
'bert.bert.encoder.layer.9.attention.self.key.weight',
'bert.bert.encoder.layer.9.attention.self.key.bias',
'bert.bert.encoder.layer.9.attention.self.value.weight',
'bert.bert.encoder.layer.9.attention.self.value.bias',
'bert.bert.encoder.layer.9.attention.output.dense.weight',
'bert.bert.encoder.layer.9.attention.output.dense.bias',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.intermediate.dense.weight',
'bert.bert.encoder.layer.9.intermediate.dense.bias',
'bert.bert.encoder.layer.9.output.dense.weight',
'bert.bert.encoder.layer.9.output.dense.bias',
'bert.bert.encoder.layer.9.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.attention.self.query.weight',
'bert.bert.encoder.layer.10.attention.self.query.bias',
'bert.bert.encoder.layer.10.attention.self.key.weight',
'bert.bert.encoder.layer.10.attention.self.key.bias',
'bert.bert.encoder.layer.10.attention.self.value.weight',

```

'bert.bert.encoder.layer.10.attention.self.value.bias',
'bert.bert.encoder.layer.10.attention.output.dense.weight',
'bert.bert.encoder.layer.10.attention.output.dense.bias',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.intermediate.dense.weight',
'bert.bert.encoder.layer.10.intermediate.dense.bias',
'bert.bert.encoder.layer.10.output.dense.weight',
'bert.bert.encoder.layer.10.output.dense.bias',
'bert.bert.encoder.layer.10.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.attention.self.query.weight',
'bert.bert.encoder.layer.11.attention.self.query.bias',
'bert.bert.encoder.layer.11.attention.self.key.weight',
'bert.bert.encoder.layer.11.attention.self.key.bias',
'bert.bert.encoder.layer.11.attention.self.value.weight',
'bert.bert.encoder.layer.11.attention.self.value.bias',
'bert.bert.encoder.layer.11.attention.output.dense.weight',
'bert.bert.encoder.layer.11.attention.output.dense.bias',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.intermediate.dense.weight',
'bert.bert.encoder.layer.11.intermediate.dense.bias',
'bert.bert.encoder.layer.11.output.dense.weight',
'bert.bert.encoder.layer.11.output.dense.bias',
'bert.bert.encoder.layer.11.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.output.LayerNorm.bias', 'bert.classifier.weight',
'bert.classifier.bias']

```

100%| | 1054/1054 [01:05<00:00, 16.18it/s]

Epochs: 4 | Train Loss: 0.023 | Train Accuracy: 0.993 | Val Loss: 0.185 | Val Accuracy: 0.956

Updated parameters during this epoch:

```

Parameters updated: ['bert.bert.embeddings.word_embeddings.weight',
'bert.bert.embeddings.position_embeddings.weight',
'bert.bert.embeddings.token_type_embeddings.weight',
'bert.bert.embeddings.LayerNorm.weight', 'bert.bert.embeddings.LayerNorm.bias',
'bert.bert.encoder.layer.0.attention.self.query.weight',
'bert.bert.encoder.layer.0.attention.self.query.bias',
'bert.bert.encoder.layer.0.attention.self.key.weight',
'bert.bert.encoder.layer.0.attention.self.key.bias',
'bert.bert.encoder.layer.0.attention.self.value.weight',
'bert.bert.encoder.layer.0.attention.self.value.bias',
'bert.bert.encoder.layer.0.attention.output.dense.weight',
'bert.bert.encoder.layer.0.attention.output.dense.bias',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.0.intermediate.dense.weight',

```


'bert.bert.encoder.layer.0.intermediate.dense.bias',
'bert.bert.encoder.layer.0.output.dense.weight',
'bert.bert.encoder.layer.0.output.dense.bias',
'bert.bert.encoder.layer.0.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.attention.self.query.weight',
'bert.bert.encoder.layer.1.attention.self.query.bias',
'bert.bert.encoder.layer.1.attention.self.key.weight',
'bert.bert.encoder.layer.1.attention.self.key.bias',
'bert.bert.encoder.layer.1.attention.self.value.weight',
'bert.bert.encoder.layer.1.attention.self.value.bias',
'bert.bert.encoder.layer.1.attention.output.dense.weight',
'bert.bert.encoder.layer.1.attention.output.dense.bias',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.intermediate.dense.weight',
'bert.bert.encoder.layer.1.intermediate.dense.bias',
'bert.bert.encoder.layer.1.output.dense.weight',
'bert.bert.encoder.layer.1.output.dense.bias',
'bert.bert.encoder.layer.1.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.attention.self.query.weight',
'bert.bert.encoder.layer.2.attention.self.query.bias',
'bert.bert.encoder.layer.2.attention.self.key.weight',
'bert.bert.encoder.layer.2.attention.self.key.bias',
'bert.bert.encoder.layer.2.attention.self.value.weight',
'bert.bert.encoder.layer.2.attention.self.value.bias',
'bert.bert.encoder.layer.2.attention.output.dense.weight',
'bert.bert.encoder.layer.2.attention.output.dense.bias',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.intermediate.dense.weight',
'bert.bert.encoder.layer.2.intermediate.dense.bias',
'bert.bert.encoder.layer.2.output.dense.weight',
'bert.bert.encoder.layer.2.output.dense.bias',
'bert.bert.encoder.layer.2.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.attention.self.query.weight',
'bert.bert.encoder.layer.3.attention.self.query.bias',
'bert.bert.encoder.layer.3.attention.self.key.weight',
'bert.bert.encoder.layer.3.attention.self.key.bias',
'bert.bert.encoder.layer.3.attention.self.value.weight',
'bert.bert.encoder.layer.3.attention.self.value.bias',
'bert.bert.encoder.layer.3.attention.output.dense.weight',
'bert.bert.encoder.layer.3.attention.output.dense.bias',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.intermediate.dense.weight',

'bert.bert.encoder.layer.3.intermediate.dense.bias',
'bert.bert.encoder.layer.3.output.dense.weight',
'bert.bert.encoder.layer.3.output.dense.bias',
'bert.bert.encoder.layer.3.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.attention.self.query.weight',
'bert.bert.encoder.layer.4.attention.self.query.bias',
'bert.bert.encoder.layer.4.attention.self.key.weight',
'bert.bert.encoder.layer.4.attention.self.key.bias',
'bert.bert.encoder.layer.4.attention.self.value.weight',
'bert.bert.encoder.layer.4.attention.self.value.bias',
'bert.bert.encoder.layer.4.attention.output.dense.weight',
'bert.bert.encoder.layer.4.attention.output.dense.bias',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.intermediate.dense.weight',
'bert.bert.encoder.layer.4.intermediate.dense.bias',
'bert.bert.encoder.layer.4.output.dense.weight',
'bert.bert.encoder.layer.4.output.dense.bias',
'bert.bert.encoder.layer.4.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.attention.self.query.weight',
'bert.bert.encoder.layer.5.attention.self.query.bias',
'bert.bert.encoder.layer.5.attention.self.key.weight',
'bert.bert.encoder.layer.5.attention.self.key.bias',
'bert.bert.encoder.layer.5.attention.self.value.weight',
'bert.bert.encoder.layer.5.attention.self.value.bias',
'bert.bert.encoder.layer.5.attention.output.dense.weight',
'bert.bert.encoder.layer.5.attention.output.dense.bias',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.intermediate.dense.weight',
'bert.bert.encoder.layer.5.intermediate.dense.bias',
'bert.bert.encoder.layer.5.output.dense.weight',
'bert.bert.encoder.layer.5.output.dense.bias',
'bert.bert.encoder.layer.5.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.attention.self.query.weight',
'bert.bert.encoder.layer.6.attention.self.query.bias',
'bert.bert.encoder.layer.6.attention.self.key.weight',
'bert.bert.encoder.layer.6.attention.self.key.bias',
'bert.bert.encoder.layer.6.attention.self.value.weight',
'bert.bert.encoder.layer.6.attention.self.value.bias',
'bert.bert.encoder.layer.6.attention.output.dense.weight',
'bert.bert.encoder.layer.6.attention.output.dense.bias',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.intermediate.dense.weight',

'bert.bert.encoder.layer.6.intermediate.dense.bias',
'bert.bert.encoder.layer.6.output.dense.weight',
'bert.bert.encoder.layer.6.output.dense.bias',
'bert.bert.encoder.layer.6.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.attention.self.query.weight',
'bert.bert.encoder.layer.7.attention.self.query.bias',
'bert.bert.encoder.layer.7.attention.self.key.weight',
'bert.bert.encoder.layer.7.attention.self.key.bias',
'bert.bert.encoder.layer.7.attention.self.value.weight',
'bert.bert.encoder.layer.7.attention.self.value.bias',
'bert.bert.encoder.layer.7.attention.output.dense.weight',
'bert.bert.encoder.layer.7.attention.output.dense.bias',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.intermediate.dense.weight',
'bert.bert.encoder.layer.7.intermediate.dense.bias',
'bert.bert.encoder.layer.7.output.dense.weight',
'bert.bert.encoder.layer.7.output.dense.bias',
'bert.bert.encoder.layer.7.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.attention.self.query.weight',
'bert.bert.encoder.layer.8.attention.self.query.bias',
'bert.bert.encoder.layer.8.attention.self.key.weight',
'bert.bert.encoder.layer.8.attention.self.key.bias',
'bert.bert.encoder.layer.8.attention.self.value.weight',
'bert.bert.encoder.layer.8.attention.self.value.bias',
'bert.bert.encoder.layer.8.attention.output.dense.weight',
'bert.bert.encoder.layer.8.attention.output.dense.bias',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.intermediate.dense.weight',
'bert.bert.encoder.layer.8.intermediate.dense.bias',
'bert.bert.encoder.layer.8.output.dense.weight',
'bert.bert.encoder.layer.8.output.dense.bias',
'bert.bert.encoder.layer.8.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.attention.self.query.weight',
'bert.bert.encoder.layer.9.attention.self.query.bias',
'bert.bert.encoder.layer.9.attention.self.key.weight',
'bert.bert.encoder.layer.9.attention.self.key.bias',
'bert.bert.encoder.layer.9.attention.self.value.weight',
'bert.bert.encoder.layer.9.attention.self.value.bias',
'bert.bert.encoder.layer.9.attention.output.dense.weight',
'bert.bert.encoder.layer.9.attention.output.dense.bias',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.intermediate.dense.weight',

```
'bert.bert.encoder.layer.9.intermediate.dense.bias',
'bert.bert.encoder.layer.9.output.dense.weight',
'bert.bert.encoder.layer.9.output.dense.bias',
'bert.bert.encoder.layer.9.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.attention.self.query.weight',
'bert.bert.encoder.layer.10.attention.self.query.bias',
'bert.bert.encoder.layer.10.attention.self.key.weight',
'bert.bert.encoder.layer.10.attention.self.key.bias',
'bert.bert.encoder.layer.10.attention.self.value.weight',
'bert.bert.encoder.layer.10.attention.self.value.bias',
'bert.bert.encoder.layer.10.attention.output.dense.weight',
'bert.bert.encoder.layer.10.attention.output.dense.bias',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.intermediate.dense.weight',
'bert.bert.encoder.layer.10.intermediate.dense.bias',
'bert.bert.encoder.layer.10.output.dense.weight',
'bert.bert.encoder.layer.10.output.dense.bias',
'bert.bert.encoder.layer.10.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.attention.self.query.weight',
'bert.bert.encoder.layer.11.attention.self.query.bias',
'bert.bert.encoder.layer.11.attention.self.key.weight',
'bert.bert.encoder.layer.11.attention.self.key.bias',
'bert.bert.encoder.layer.11.attention.self.value.weight',
'bert.bert.encoder.layer.11.attention.self.value.bias',
'bert.bert.encoder.layer.11.attention.output.dense.weight',
'bert.bert.encoder.layer.11.attention.output.dense.bias',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.intermediate.dense.weight',
'bert.bert.encoder.layer.11.intermediate.dense.bias',
'bert.bert.encoder.layer.11.output.dense.weight',
'bert.bert.encoder.layer.11.output.dense.bias',
'bert.bert.encoder.layer.11.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.output.LayerNorm.bias', 'bert.classifier.weight',
'bert.classifier.bias']
```

100%| | 1054/1054 [01:05<00:00, 16.17it/s]

Epochs: 5 | Train Loss: 0.005 | Train Accuracy: 0.999 | Val Loss: 0.228 | Val Accuracy: 0.953

Updated parameters during this epoch:

```
Parameters updated: ['bert.bert.embeddings.word_embeddings.weight',
'bert.bert.embeddings.position_embeddings.weight',
'bert.bert.embeddings.token_type_embeddings.weight',
'bert.bert.embeddings.LayerNorm.weight', 'bert.bert.embeddings.LayerNorm.bias',
'bert.bert.encoder.layer.0.attention.self.query.weight',
```

'bert.bert.encoder.layer.0.attention.self.query.bias',
'bert.bert.encoder.layer.0.attention.self.key.weight',
'bert.bert.encoder.layer.0.attention.self.key.bias',
'bert.bert.encoder.layer.0.attention.self.value.weight',
'bert.bert.encoder.layer.0.attention.self.value.bias',
'bert.bert.encoder.layer.0.attention.output.dense.weight',
'bert.bert.encoder.layer.0.attention.output.dense.bias',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.0.intermediate.dense.weight',
'bert.bert.encoder.layer.0.intermediate.dense.bias',
'bert.bert.encoder.layer.0.output.dense.weight',
'bert.bert.encoder.layer.0.output.dense.bias',
'bert.bert.encoder.layer.0.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.attention.self.query.weight',
'bert.bert.encoder.layer.1.attention.self.query.bias',
'bert.bert.encoder.layer.1.attention.self.key.weight',
'bert.bert.encoder.layer.1.attention.self.key.bias',
'bert.bert.encoder.layer.1.attention.self.value.weight',
'bert.bert.encoder.layer.1.attention.self.value.bias',
'bert.bert.encoder.layer.1.attention.output.dense.weight',
'bert.bert.encoder.layer.1.attention.output.dense.bias',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.intermediate.dense.weight',
'bert.bert.encoder.layer.1.intermediate.dense.bias',
'bert.bert.encoder.layer.1.output.dense.weight',
'bert.bert.encoder.layer.1.output.dense.bias',
'bert.bert.encoder.layer.1.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.attention.self.query.weight',
'bert.bert.encoder.layer.2.attention.self.query.bias',
'bert.bert.encoder.layer.2.attention.self.key.weight',
'bert.bert.encoder.layer.2.attention.self.key.bias',
'bert.bert.encoder.layer.2.attention.self.value.weight',
'bert.bert.encoder.layer.2.attention.self.value.bias',
'bert.bert.encoder.layer.2.attention.output.dense.weight',
'bert.bert.encoder.layer.2.attention.output.dense.bias',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.intermediate.dense.weight',
'bert.bert.encoder.layer.2.intermediate.dense.bias',
'bert.bert.encoder.layer.2.output.dense.weight',
'bert.bert.encoder.layer.2.output.dense.bias',
'bert.bert.encoder.layer.2.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.attention.self.query.weight',

'bert.bert.encoder.layer.3.attention.self.query.bias',
'bert.bert.encoder.layer.3.attention.self.key.weight',
'bert.bert.encoder.layer.3.attention.self.key.bias',
'bert.bert.encoder.layer.3.attention.self.value.weight',
'bert.bert.encoder.layer.3.attention.self.value.bias',
'bert.bert.encoder.layer.3.attention.output.dense.weight',
'bert.bert.encoder.layer.3.attention.output.dense.bias',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.intermediate.dense.weight',
'bert.bert.encoder.layer.3.intermediate.dense.bias',
'bert.bert.encoder.layer.3.output.dense.weight',
'bert.bert.encoder.layer.3.output.dense.bias',
'bert.bert.encoder.layer.3.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.attention.self.query.weight',
'bert.bert.encoder.layer.4.attention.self.query.bias',
'bert.bert.encoder.layer.4.attention.self.key.weight',
'bert.bert.encoder.layer.4.attention.self.key.bias',
'bert.bert.encoder.layer.4.attention.self.value.weight',
'bert.bert.encoder.layer.4.attention.self.value.bias',
'bert.bert.encoder.layer.4.attention.output.dense.weight',
'bert.bert.encoder.layer.4.attention.output.dense.bias',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.intermediate.dense.weight',
'bert.bert.encoder.layer.4.intermediate.dense.bias',
'bert.bert.encoder.layer.4.output.dense.weight',
'bert.bert.encoder.layer.4.output.dense.bias',
'bert.bert.encoder.layer.4.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.attention.self.query.weight',
'bert.bert.encoder.layer.5.attention.self.query.bias',
'bert.bert.encoder.layer.5.attention.self.key.weight',
'bert.bert.encoder.layer.5.attention.self.key.bias',
'bert.bert.encoder.layer.5.attention.self.value.weight',
'bert.bert.encoder.layer.5.attention.self.value.bias',
'bert.bert.encoder.layer.5.attention.output.dense.weight',
'bert.bert.encoder.layer.5.attention.output.dense.bias',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.intermediate.dense.weight',
'bert.bert.encoder.layer.5.intermediate.dense.bias',
'bert.bert.encoder.layer.5.output.dense.weight',
'bert.bert.encoder.layer.5.output.dense.bias',
'bert.bert.encoder.layer.5.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.attention.self.query.weight',

'bert.bert.encoder.layer.6.attention.self.query.bias',
'bert.bert.encoder.layer.6.attention.self.key.weight',
'bert.bert.encoder.layer.6.attention.self.key.bias',
'bert.bert.encoder.layer.6.attention.self.value.weight',
'bert.bert.encoder.layer.6.attention.self.value.bias',
'bert.bert.encoder.layer.6.attention.output.dense.weight',
'bert.bert.encoder.layer.6.attention.output.dense.bias',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.intermediate.dense.weight',
'bert.bert.encoder.layer.6.intermediate.dense.bias',
'bert.bert.encoder.layer.6.output.dense.weight',
'bert.bert.encoder.layer.6.output.dense.bias',
'bert.bert.encoder.layer.6.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.attention.self.query.weight',
'bert.bert.encoder.layer.7.attention.self.query.bias',
'bert.bert.encoder.layer.7.attention.self.key.weight',
'bert.bert.encoder.layer.7.attention.self.key.bias',
'bert.bert.encoder.layer.7.attention.self.value.weight',
'bert.bert.encoder.layer.7.attention.self.value.bias',
'bert.bert.encoder.layer.7.attention.output.dense.weight',
'bert.bert.encoder.layer.7.attention.output.dense.bias',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.intermediate.dense.weight',
'bert.bert.encoder.layer.7.intermediate.dense.bias',
'bert.bert.encoder.layer.7.output.dense.weight',
'bert.bert.encoder.layer.7.output.dense.bias',
'bert.bert.encoder.layer.7.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.attention.self.query.weight',
'bert.bert.encoder.layer.8.attention.self.query.bias',
'bert.bert.encoder.layer.8.attention.self.key.weight',
'bert.bert.encoder.layer.8.attention.self.key.bias',
'bert.bert.encoder.layer.8.attention.self.value.weight',
'bert.bert.encoder.layer.8.attention.self.value.bias',
'bert.bert.encoder.layer.8.attention.output.dense.weight',
'bert.bert.encoder.layer.8.attention.output.dense.bias',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.intermediate.dense.weight',
'bert.bert.encoder.layer.8.intermediate.dense.bias',
'bert.bert.encoder.layer.8.output.dense.weight',
'bert.bert.encoder.layer.8.output.dense.bias',
'bert.bert.encoder.layer.8.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.attention.self.query.weight',

```

'bert.bert.encoder.layer.9.attention.self.query.bias',
'bert.bert.encoder.layer.9.attention.self.key.weight',
'bert.bert.encoder.layer.9.attention.self.key.bias',
'bert.bert.encoder.layer.9.attention.self.value.weight',
'bert.bert.encoder.layer.9.attention.self.value.bias',
'bert.bert.encoder.layer.9.attention.output.dense.weight',
'bert.bert.encoder.layer.9.attention.output.dense.bias',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.intermediate.dense.weight',
'bert.bert.encoder.layer.9.intermediate.dense.bias',
'bert.bert.encoder.layer.9.output.dense.weight',
'bert.bert.encoder.layer.9.output.dense.bias',
'bert.bert.encoder.layer.9.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.attention.self.query.weight',
'bert.bert.encoder.layer.10.attention.self.query.bias',
'bert.bert.encoder.layer.10.attention.self.key.weight',
'bert.bert.encoder.layer.10.attention.self.key.bias',
'bert.bert.encoder.layer.10.attention.self.value.weight',
'bert.bert.encoder.layer.10.attention.self.value.bias',
'bert.bert.encoder.layer.10.attention.output.dense.weight',
'bert.bert.encoder.layer.10.attention.output.dense.bias',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.intermediate.dense.weight',
'bert.bert.encoder.layer.10.intermediate.dense.bias',
'bert.bert.encoder.layer.10.output.dense.weight',
'bert.bert.encoder.layer.10.output.dense.bias',
'bert.bert.encoder.layer.10.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.attention.self.query.weight',
'bert.bert.encoder.layer.11.attention.self.query.bias',
'bert.bert.encoder.layer.11.attention.self.key.weight',
'bert.bert.encoder.layer.11.attention.self.key.bias',
'bert.bert.encoder.layer.11.attention.self.value.weight',
'bert.bert.encoder.layer.11.attention.self.value.bias',
'bert.bert.encoder.layer.11.attention.output.dense.weight',
'bert.bert.encoder.layer.11.attention.output.dense.bias',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.intermediate.dense.weight',
'bert.bert.encoder.layer.11.intermediate.dense.bias',
'bert.bert.encoder.layer.11.output.dense.weight',
'bert.bert.encoder.layer.11.output.dense.bias',
'bert.bert.encoder.layer.11.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.output.LayerNorm.bias', 'bert.classifier.weight',
'bert.classifier.bias']

```


100%| | 1054/1054 [01:05<00:00, 16.18it/s]

Epochs: 6 | Train Loss: 0.002 | Train Accuracy: 1.000 | Val Loss: 0.240 | Val Accuracy: 0.953

Updated parameters during this epoch:

Parameters updated: ['bert.bert.embeddings.word_embeddings.weight',
'bert.bert.embeddings.position_embeddings.weight',
'bert.bert.embeddings.token_type_embeddings.weight',
'bert.bert.embeddings.LayerNorm.weight', 'bert.bert.embeddings.LayerNorm.bias',
'bert.bert.encoder.layer.0.attention.self.query.weight',
'bert.bert.encoder.layer.0.attention.self.query.bias',
'bert.bert.encoder.layer.0.attention.self.key.weight',
'bert.bert.encoder.layer.0.attention.self.key.bias',
'bert.bert.encoder.layer.0.attention.self.value.weight',
'bert.bert.encoder.layer.0.attention.self.value.bias',
'bert.bert.encoder.layer.0.attention.output.dense.weight',
'bert.bert.encoder.layer.0.attention.output.dense.bias',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.0.intermediate.dense.weight',
'bert.bert.encoder.layer.0.intermediate.dense.bias',
'bert.bert.encoder.layer.0.output.dense.weight',
'bert.bert.encoder.layer.0.output.dense.bias',
'bert.bert.encoder.layer.0.output.LayerNorm.weight',
'bert.bert.encoder.layer.0.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.attention.self.query.weight',
'bert.bert.encoder.layer.1.attention.self.query.bias',
'bert.bert.encoder.layer.1.attention.self.key.weight',
'bert.bert.encoder.layer.1.attention.self.key.bias',
'bert.bert.encoder.layer.1.attention.self.value.weight',
'bert.bert.encoder.layer.1.attention.self.value.bias',
'bert.bert.encoder.layer.1.attention.output.dense.weight',
'bert.bert.encoder.layer.1.attention.output.dense.bias',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.1.intermediate.dense.weight',
'bert.bert.encoder.layer.1.intermediate.dense.bias',
'bert.bert.encoder.layer.1.output.dense.weight',
'bert.bert.encoder.layer.1.output.dense.bias',
'bert.bert.encoder.layer.1.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.attention.self.query.weight',
'bert.bert.encoder.layer.2.attention.self.query.bias',
'bert.bert.encoder.layer.2.attention.self.key.weight',
'bert.bert.encoder.layer.2.attention.self.key.bias',
'bert.bert.encoder.layer.2.attention.self.value.weight',
'bert.bert.encoder.layer.2.attention.self.value.bias',
'bert.bert.encoder.layer.2.attention.output.dense.weight',

'bert.bert.encoder.layer.2.attention.output.dense.bias',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.intermediate.dense.weight',
'bert.bert.encoder.layer.2.intermediate.dense.bias',
'bert.bert.encoder.layer.2.output.dense.weight',
'bert.bert.encoder.layer.2.output.dense.bias',
'bert.bert.encoder.layer.2.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.attention.self.query.weight',
'bert.bert.encoder.layer.3.attention.self.query.bias',
'bert.bert.encoder.layer.3.attention.self.key.weight',
'bert.bert.encoder.layer.3.attention.self.key.bias',
'bert.bert.encoder.layer.3.attention.self.value.weight',
'bert.bert.encoder.layer.3.attention.self.value.bias',
'bert.bert.encoder.layer.3.attention.output.dense.weight',
'bert.bert.encoder.layer.3.attention.output.dense.bias',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.intermediate.dense.weight',
'bert.bert.encoder.layer.3.intermediate.dense.bias',
'bert.bert.encoder.layer.3.output.dense.weight',
'bert.bert.encoder.layer.3.output.dense.bias',
'bert.bert.encoder.layer.3.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.attention.self.query.weight',
'bert.bert.encoder.layer.4.attention.self.query.bias',
'bert.bert.encoder.layer.4.attention.self.key.weight',
'bert.bert.encoder.layer.4.attention.self.key.bias',
'bert.bert.encoder.layer.4.attention.self.value.weight',
'bert.bert.encoder.layer.4.attention.self.value.bias',
'bert.bert.encoder.layer.4.attention.output.dense.weight',
'bert.bert.encoder.layer.4.attention.output.dense.bias',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.intermediate.dense.weight',
'bert.bert.encoder.layer.4.intermediate.dense.bias',
'bert.bert.encoder.layer.4.output.dense.weight',
'bert.bert.encoder.layer.4.output.dense.bias',
'bert.bert.encoder.layer.4.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.attention.self.query.weight',
'bert.bert.encoder.layer.5.attention.self.query.bias',
'bert.bert.encoder.layer.5.attention.self.key.weight',
'bert.bert.encoder.layer.5.attention.self.key.bias',
'bert.bert.encoder.layer.5.attention.self.value.weight',
'bert.bert.encoder.layer.5.attention.self.value.bias',
'bert.bert.encoder.layer.5.attention.output.dense.weight',

'bert.bert.encoder.layer.5.attention.output.dense.bias',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.intermediate.dense.weight',
'bert.bert.encoder.layer.5.intermediate.dense.bias',
'bert.bert.encoder.layer.5.output.dense.weight',
'bert.bert.encoder.layer.5.output.dense.bias',
'bert.bert.encoder.layer.5.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.attention.self.query.weight',
'bert.bert.encoder.layer.6.attention.self.query.bias',
'bert.bert.encoder.layer.6.attention.self.key.weight',
'bert.bert.encoder.layer.6.attention.self.key.bias',
'bert.bert.encoder.layer.6.attention.self.value.weight',
'bert.bert.encoder.layer.6.attention.self.value.bias',
'bert.bert.encoder.layer.6.attention.output.dense.weight',
'bert.bert.encoder.layer.6.attention.output.dense.bias',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.intermediate.dense.weight',
'bert.bert.encoder.layer.6.intermediate.dense.bias',
'bert.bert.encoder.layer.6.output.dense.weight',
'bert.bert.encoder.layer.6.output.dense.bias',
'bert.bert.encoder.layer.6.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.attention.self.query.weight',
'bert.bert.encoder.layer.7.attention.self.query.bias',
'bert.bert.encoder.layer.7.attention.self.key.weight',
'bert.bert.encoder.layer.7.attention.self.key.bias',
'bert.bert.encoder.layer.7.attention.self.value.weight',
'bert.bert.encoder.layer.7.attention.self.value.bias',
'bert.bert.encoder.layer.7.attention.output.dense.weight',
'bert.bert.encoder.layer.7.attention.output.dense.bias',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.intermediate.dense.weight',
'bert.bert.encoder.layer.7.intermediate.dense.bias',
'bert.bert.encoder.layer.7.output.dense.weight',
'bert.bert.encoder.layer.7.output.dense.bias',
'bert.bert.encoder.layer.7.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.attention.self.query.weight',
'bert.bert.encoder.layer.8.attention.self.query.bias',
'bert.bert.encoder.layer.8.attention.self.key.weight',
'bert.bert.encoder.layer.8.attention.self.key.bias',
'bert.bert.encoder.layer.8.attention.self.value.weight',
'bert.bert.encoder.layer.8.attention.self.value.bias',
'bert.bert.encoder.layer.8.attention.output.dense.weight',

'bert.bert.encoder.layer.8.attention.output.dense.bias',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.intermediate.dense.weight',
'bert.bert.encoder.layer.8.intermediate.dense.bias',
'bert.bert.encoder.layer.8.output.dense.weight',
'bert.bert.encoder.layer.8.output.dense.bias',
'bert.bert.encoder.layer.8.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.attention.self.query.weight',
'bert.bert.encoder.layer.9.attention.self.query.bias',
'bert.bert.encoder.layer.9.attention.self.key.weight',
'bert.bert.encoder.layer.9.attention.self.key.bias',
'bert.bert.encoder.layer.9.attention.self.value.weight',
'bert.bert.encoder.layer.9.attention.self.value.bias',
'bert.bert.encoder.layer.9.attention.output.dense.weight',
'bert.bert.encoder.layer.9.attention.output.dense.bias',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.intermediate.dense.weight',
'bert.bert.encoder.layer.9.intermediate.dense.bias',
'bert.bert.encoder.layer.9.output.dense.weight',
'bert.bert.encoder.layer.9.output.dense.bias',
'bert.bert.encoder.layer.9.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.attention.self.query.weight',
'bert.bert.encoder.layer.10.attention.self.query.bias',
'bert.bert.encoder.layer.10.attention.self.key.weight',
'bert.bert.encoder.layer.10.attention.self.key.bias',
'bert.bert.encoder.layer.10.attention.self.value.weight',
'bert.bert.encoder.layer.10.attention.self.value.bias',
'bert.bert.encoder.layer.10.attention.output.dense.weight',
'bert.bert.encoder.layer.10.attention.output.dense.bias',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.intermediate.dense.weight',
'bert.bert.encoder.layer.10.intermediate.dense.bias',
'bert.bert.encoder.layer.10.output.dense.weight',
'bert.bert.encoder.layer.10.output.dense.bias',
'bert.bert.encoder.layer.10.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.attention.self.query.weight',
'bert.bert.encoder.layer.11.attention.self.query.bias',
'bert.bert.encoder.layer.11.attention.self.key.weight',
'bert.bert.encoder.layer.11.attention.self.key.bias',
'bert.bert.encoder.layer.11.attention.self.value.weight',
'bert.bert.encoder.layer.11.attention.self.value.bias',
'bert.bert.encoder.layer.11.attention.output.dense.weight',

```
'bert.bert.encoder.layer.11.attention.output.dense.bias',  
'bert.bert.encoder.layer.11.attention.output.LayerNorm.weight',  
'bert.bert.encoder.layer.11.attention.output.LayerNorm.bias',  
'bert.bert.encoder.layer.11.intermediate.dense.weight',  
'bert.bert.encoder.layer.11.intermediate.dense.bias',  
'bert.bert.encoder.layer.11.output.dense.weight',  
'bert.bert.encoder.layer.11.output.dense.bias',  
'bert.bert.encoder.layer.11.output.LayerNorm.weight',  
'bert.bert.encoder.layer.11.output.LayerNorm.bias', 'bert.classifier.weight',  
'bert.classifier.bias']
```

```
100%|      | 1054/1054 [01:05<00:00, 16.16it/s]
```

```
Epochs: 7 | Train Loss: 0.001 | Train Accuracy: 1.000 | Val Loss: 0.251 | Val  
Accuracy: 0.953
```

```
Updated parameters during this epoch:
```

```
Parameters updated: ['bert.bert.embeddings.word_embeddings.weight',  
'bert.bert.embeddings.position_embeddings.weight',  
'bert.bert.embeddings.token_type_embeddings.weight',  
'bert.bert.embeddings.LayerNorm.weight', 'bert.bert.embeddings.LayerNorm.bias',  
'bert.bert.encoder.layer.0.attention.self.query.weight',  
'bert.bert.encoder.layer.0.attention.self.query.bias',  
'bert.bert.encoder.layer.0.attention.self.key.weight',  
'bert.bert.encoder.layer.0.attention.self.key.bias',  
'bert.bert.encoder.layer.0.attention.self.value.weight',  
'bert.bert.encoder.layer.0.attention.self.value.bias',  
'bert.bert.encoder.layer.0.attention.output.dense.weight',  
'bert.bert.encoder.layer.0.attention.output.dense.bias',  
'bert.bert.encoder.layer.0.attention.output.LayerNorm.weight',  
'bert.bert.encoder.layer.0.attention.output.LayerNorm.bias',  
'bert.bert.encoder.layer.0.intermediate.dense.weight',  
'bert.bert.encoder.layer.0.intermediate.dense.bias',  
'bert.bert.encoder.layer.0.output.dense.weight',  
'bert.bert.encoder.layer.0.output.dense.bias',  
'bert.bert.encoder.layer.0.output.LayerNorm.weight',  
'bert.bert.encoder.layer.0.output.LayerNorm.bias',  
'bert.bert.encoder.layer.1.attention.self.query.weight',  
'bert.bert.encoder.layer.1.attention.self.query.bias',  
'bert.bert.encoder.layer.1.attention.self.key.weight',  
'bert.bert.encoder.layer.1.attention.self.key.bias',  
'bert.bert.encoder.layer.1.attention.self.value.weight',  
'bert.bert.encoder.layer.1.attention.self.value.bias',  
'bert.bert.encoder.layer.1.attention.output.dense.weight',  
'bert.bert.encoder.layer.1.attention.output.dense.bias',  
'bert.bert.encoder.layer.1.attention.output.LayerNorm.weight',  
'bert.bert.encoder.layer.1.attention.output.LayerNorm.bias',  
'bert.bert.encoder.layer.1.intermediate.dense.weight',  
'bert.bert.encoder.layer.1.intermediate.dense.bias',  
'bert.bert.encoder.layer.1.output.dense.weight',
```

'bert.bert.encoder.layer.1.output.dense.bias',
'bert.bert.encoder.layer.1.output.LayerNorm.weight',
'bert.bert.encoder.layer.1.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.attention.self.query.weight',
'bert.bert.encoder.layer.2.attention.self.query.bias',
'bert.bert.encoder.layer.2.attention.self.key.weight',
'bert.bert.encoder.layer.2.attention.self.key.bias',
'bert.bert.encoder.layer.2.attention.self.value.weight',
'bert.bert.encoder.layer.2.attention.self.value.bias',
'bert.bert.encoder.layer.2.attention.output.dense.weight',
'bert.bert.encoder.layer.2.attention.output.dense.bias',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.2.intermediate.dense.weight',
'bert.bert.encoder.layer.2.intermediate.dense.bias',
'bert.bert.encoder.layer.2.output.dense.weight',
'bert.bert.encoder.layer.2.output.dense.bias',
'bert.bert.encoder.layer.2.output.LayerNorm.weight',
'bert.bert.encoder.layer.2.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.attention.self.query.weight',
'bert.bert.encoder.layer.3.attention.self.query.bias',
'bert.bert.encoder.layer.3.attention.self.key.weight',
'bert.bert.encoder.layer.3.attention.self.key.bias',
'bert.bert.encoder.layer.3.attention.self.value.weight',
'bert.bert.encoder.layer.3.attention.self.value.bias',
'bert.bert.encoder.layer.3.attention.output.dense.weight',
'bert.bert.encoder.layer.3.attention.output.dense.bias',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.3.intermediate.dense.weight',
'bert.bert.encoder.layer.3.intermediate.dense.bias',
'bert.bert.encoder.layer.3.output.dense.weight',
'bert.bert.encoder.layer.3.output.dense.bias',
'bert.bert.encoder.layer.3.output.LayerNorm.weight',
'bert.bert.encoder.layer.3.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.attention.self.query.weight',
'bert.bert.encoder.layer.4.attention.self.query.bias',
'bert.bert.encoder.layer.4.attention.self.key.weight',
'bert.bert.encoder.layer.4.attention.self.key.bias',
'bert.bert.encoder.layer.4.attention.self.value.weight',
'bert.bert.encoder.layer.4.attention.self.value.bias',
'bert.bert.encoder.layer.4.attention.output.dense.weight',
'bert.bert.encoder.layer.4.attention.output.dense.bias',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.4.intermediate.dense.weight',
'bert.bert.encoder.layer.4.intermediate.dense.bias',
'bert.bert.encoder.layer.4.output.dense.weight',

'bert.bert.encoder.layer.4.output.dense.bias',
'bert.bert.encoder.layer.4.output.LayerNorm.weight',
'bert.bert.encoder.layer.4.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.attention.self.query.weight',
'bert.bert.encoder.layer.5.attention.self.query.bias',
'bert.bert.encoder.layer.5.attention.self.key.weight',
'bert.bert.encoder.layer.5.attention.self.key.bias',
'bert.bert.encoder.layer.5.attention.self.value.weight',
'bert.bert.encoder.layer.5.attention.self.value.bias',
'bert.bert.encoder.layer.5.attention.output.dense.weight',
'bert.bert.encoder.layer.5.attention.output.dense.bias',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.5.intermediate.dense.weight',
'bert.bert.encoder.layer.5.intermediate.dense.bias',
'bert.bert.encoder.layer.5.output.dense.weight',
'bert.bert.encoder.layer.5.output.dense.bias',
'bert.bert.encoder.layer.5.output.LayerNorm.weight',
'bert.bert.encoder.layer.5.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.attention.self.query.weight',
'bert.bert.encoder.layer.6.attention.self.query.bias',
'bert.bert.encoder.layer.6.attention.self.key.weight',
'bert.bert.encoder.layer.6.attention.self.key.bias',
'bert.bert.encoder.layer.6.attention.self.value.weight',
'bert.bert.encoder.layer.6.attention.self.value.bias',
'bert.bert.encoder.layer.6.attention.output.dense.weight',
'bert.bert.encoder.layer.6.attention.output.dense.bias',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.6.intermediate.dense.weight',
'bert.bert.encoder.layer.6.intermediate.dense.bias',
'bert.bert.encoder.layer.6.output.dense.weight',
'bert.bert.encoder.layer.6.output.dense.bias',
'bert.bert.encoder.layer.6.output.LayerNorm.weight',
'bert.bert.encoder.layer.6.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.attention.self.query.weight',
'bert.bert.encoder.layer.7.attention.self.query.bias',
'bert.bert.encoder.layer.7.attention.self.key.weight',
'bert.bert.encoder.layer.7.attention.self.key.bias',
'bert.bert.encoder.layer.7.attention.self.value.weight',
'bert.bert.encoder.layer.7.attention.self.value.bias',
'bert.bert.encoder.layer.7.attention.output.dense.weight',
'bert.bert.encoder.layer.7.attention.output.dense.bias',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.7.intermediate.dense.weight',
'bert.bert.encoder.layer.7.intermediate.dense.bias',
'bert.bert.encoder.layer.7.output.dense.weight',

'bert.bert.encoder.layer.7.output.dense.bias',
'bert.bert.encoder.layer.7.output.LayerNorm.weight',
'bert.bert.encoder.layer.7.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.attention.self.query.weight',
'bert.bert.encoder.layer.8.attention.self.query.bias',
'bert.bert.encoder.layer.8.attention.self.key.weight',
'bert.bert.encoder.layer.8.attention.self.key.bias',
'bert.bert.encoder.layer.8.attention.self.value.weight',
'bert.bert.encoder.layer.8.attention.self.value.bias',
'bert.bert.encoder.layer.8.attention.output.dense.weight',
'bert.bert.encoder.layer.8.attention.output.dense.bias',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.8.intermediate.dense.weight',
'bert.bert.encoder.layer.8.intermediate.dense.bias',
'bert.bert.encoder.layer.8.output.dense.weight',
'bert.bert.encoder.layer.8.output.dense.bias',
'bert.bert.encoder.layer.8.output.LayerNorm.weight',
'bert.bert.encoder.layer.8.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.attention.self.query.weight',
'bert.bert.encoder.layer.9.attention.self.query.bias',
'bert.bert.encoder.layer.9.attention.self.key.weight',
'bert.bert.encoder.layer.9.attention.self.key.bias',
'bert.bert.encoder.layer.9.attention.self.value.weight',
'bert.bert.encoder.layer.9.attention.self.value.bias',
'bert.bert.encoder.layer.9.attention.output.dense.weight',
'bert.bert.encoder.layer.9.attention.output.dense.bias',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.9.intermediate.dense.weight',
'bert.bert.encoder.layer.9.intermediate.dense.bias',
'bert.bert.encoder.layer.9.output.dense.weight',
'bert.bert.encoder.layer.9.output.dense.bias',
'bert.bert.encoder.layer.9.output.LayerNorm.weight',
'bert.bert.encoder.layer.9.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.attention.self.query.weight',
'bert.bert.encoder.layer.10.attention.self.query.bias',
'bert.bert.encoder.layer.10.attention.self.key.weight',
'bert.bert.encoder.layer.10.attention.self.key.bias',
'bert.bert.encoder.layer.10.attention.self.value.weight',
'bert.bert.encoder.layer.10.attention.self.value.bias',
'bert.bert.encoder.layer.10.attention.output.dense.weight',
'bert.bert.encoder.layer.10.attention.output.dense.bias',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.10.intermediate.dense.weight',
'bert.bert.encoder.layer.10.intermediate.dense.bias',
'bert.bert.encoder.layer.10.output.dense.weight',


```

'bert.bert.encoder.layer.10.output.dense.bias',
'bert.bert.encoder.layer.10.output.LayerNorm.weight',
'bert.bert.encoder.layer.10.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.attention.self.query.weight',
'bert.bert.encoder.layer.11.attention.self.query.bias',
'bert.bert.encoder.layer.11.attention.self.key.weight',
'bert.bert.encoder.layer.11.attention.self.key.bias',
'bert.bert.encoder.layer.11.attention.self.value.weight',
'bert.bert.encoder.layer.11.attention.self.value.bias',
'bert.bert.encoder.layer.11.attention.output.dense.weight',
'bert.bert.encoder.layer.11.attention.output.dense.bias',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.attention.output.LayerNorm.bias',
'bert.bert.encoder.layer.11.intermediate.dense.weight',
'bert.bert.encoder.layer.11.intermediate.dense.bias',
'bert.bert.encoder.layer.11.output.dense.weight',
'bert.bert.encoder.layer.11.output.dense.bias',
'bert.bert.encoder.layer.11.output.LayerNorm.weight',
'bert.bert.encoder.layer.11.output.LayerNorm.bias', 'bert.classifier.weight',
'bert.classifier.bias']

```

Early stopping triggered after 5 epochs without improvement.

Stopping training early!

```

[10]: import torch
import numpy as np
from tqdm import tqdm # Optional, for progress bar
import torch
from torch import nn

def evaluate_testset(model, test_data, labels):
    test = Dataset(test_data, labels)
    test_dataloader = torch.utils.data.DataLoader(test, batch_size=1,
    ↪shuffle=False)

    use_cuda = torch.cuda.is_available()
    device = torch.device("cuda" if use_cuda else "cpu")
    model.to(device)

    model.eval()

    total_acc_test = 0
    cls_embeddings = []

```

```

total_samples = 0

with torch.no_grad():
    # Iterate over each batch in the test set
    for batch in tqdm(test_dataloader, desc="Evaluating"):
        # Extract batch data
        input_ids, attention_mask, test_label, document_ids = batch

        if input_ids is None: # Skip any empty batches
            continue

        # Move to device
        test_label = test_label.to(device)
        attention_mask = attention_mask.to(device)
        input_ids = input_ids.to(device)

        # Forward pass through the model
        logits = model(input_ids, attention_mask)
        #print(test_label)
        #print(logits)

        '''
        logit_class_1 = logits[0, 1]
        #print(logit_class_1)
        sigmoid = nn.Sigmoid()
        expertise_score = sigmoid(logit_class_1)
        print(test_label)
        print(logit_class_1)
        print(expertise_score)
        '''

        temperature = 3.0 # A temperature value > 1 softens the
↪ probabilities

        logits_scaled = logits / temperature
        sigmoid = nn.Sigmoid()
        prob = sigmoid(logits_scaled)
        print(test_label)
        print(prob[0, 1])
        '''

        #print(logits)
        # Get predicted class labels from logits
        predicted_labels = logits.argmax(dim=1)

```

```

        #print(predicted_labels)

        # Store the CLS embeddings along with their document_id
        for cls_embedding, doc_id in zip(model.cls_embeddings,
↪document_ids):
            # Append a tuple (cls_embedding, document_id) to the list
            cls_embeddings.append((cls_embedding, doc_id.item()))

        # Calculate accuracy for this batch (based on the predicted labels)
        correct_preds = (predicted_labels == test_label).sum().item()
        total_acc_test += correct_preds
        total_samples += test_label.size(0) # Add number of samples in the
↪current batch

        # Calculate overall accuracy
        accuracy = total_acc_test / total_samples

        # Print the overall accuracy
        print(f'Test Accuracy: {accuracy: .3f}')

        return cls_embeddings

```

```
[11]: cls_embeddings_test = evaluate_testset(model, test_df, labels)
```

Evaluating: 100%| | 261/261 [00:02<00:00, 107.34it/s]

Test Accuracy: 0.939

```
[12]: tensors_list_beginner = [torch.tensor(t[1]) for t in
↪cls_embeddings_with_text_train if t[2].item() == 0] # Convert to tensors
tensors_list_expert = [torch.tensor(t[1]) for t in
↪cls_embeddings_with_text_train if t[2].item() == 1] # Convert to tensors
mean_tensor_beginner = torch.mean(torch.stack(tensors_list_beginner), dim=0)
mean_tensor_expert = torch.mean(torch.stack(tensors_list_expert), dim=0)

```

```
[13]: # Initialize a defaultdict to accumulate vectors and count for each label
from collections import defaultdict

aggregated = defaultdict(lambda: {'sum': [], 'count': 0})

# Group by label, accumulate sum of vectors, and count the occurrences
for cls_embedding, doc_id in cls_embeddings_test:
    aggregated[doc_id]['sum'].append(cls_embedding)
    aggregated[doc_id]['count'] += 1

# Calculate the mean (average) for each group

```

```

for doc_id in aggregated:
    aggregated[doc_id]['mean'] = torch.mean(torch.
↳stack(aggregated[doc_id]['sum']), dim=0)

```

```

[14]: from sklearn.metrics.pairwise import cosine_similarity
correct = 0
beginner_correct = 0
expert_correct = 0

for doc_id in aggregated.keys() :
    test_mean_cls_embedding = aggregated[doc_id]['mean']

    if test_df['label'][doc_id] == 0 :
        truth = "beginner"

    elif test_df['label'][doc_id] == 1 :
        truth = "expert"

    cls_test_normalized = test_mean_cls_embedding.cpu().detach().numpy() / np.
↳linalg.norm(test_mean_cls_embedding.cpu().detach().numpy())

    cls_beginner_normalized = mean_tensor_beginner.cpu().detach().numpy() / np.
↳linalg.norm(mean_tensor_beginner.cpu().detach().numpy())
    similarity_beg = cosine_similarity(cls_test_normalized.reshape(1, -1),
↳cls_beginner_normalized.reshape(1, -1))[0][0]

    cls_expert_normalized = mean_tensor_expert.cpu().detach().numpy() / np.linalg.
↳norm(mean_tensor_expert.cpu().detach().numpy())
    similarity_expert = cosine_similarity(cls_test_normalized.reshape(1, -1),
↳cls_expert_normalized.reshape(1, -1))[0][0]

    if similarity_beg > similarity_expert and truth == "beginner" :
        beginner_correct += 1
        correct += 1
    elif similarity_beg < similarity_expert and truth == "expert" :
        expert_correct += 1
        correct += 1

print("expert correct : ", expert_correct/len(test_df[test_df['label'] ==
↳1]), "\n")
print("beginner correct : ", beginner_correct/len(test_df[test_df['label'] ==
↳0]), "\n")
print("correct : ", correct/len(test_df), "\n")

```

expert correct : 0.9787234042553191

beginner correct : 0.9454545454545454

correct : 0.9607843137254902

```
[15]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.metrics.pairwise import cosine_similarity

# Initialize counters for the confusion matrix components
tp_beg = 0 # True positives for beginner
fp_beg = 0 # False positives for beginner
fn_beg = 0 # False negatives for beginner
tp_exp = 0 # True positives for expert
fp_exp = 0 # False positives for expert
fn_exp = 0 # False negatives for expert

# Iterate over the aggregated dictionary (you can adjust based on your data
↳structure)
for doc_id in aggregated.keys():
    test_cls_embedding = aggregated[doc_id]['mean']

    # Determine the ground truth label (beginner or expert)
    truth = "beginner" if test_df['label'][doc_id] == 0 else "expert"

    # Normalize vectors
    cls_test_normalized = test_cls_embedding.cpu().detach().numpy() / np.linalg.
↳norm(test_cls_embedding.cpu().detach().numpy())
    cls_normalized_beg = mean_tensor_beginner.cpu().detach().numpy() / np.
↳linalg.norm(mean_tensor_beginner.cpu().detach().numpy())
    cls_normalized_exp = mean_tensor_expert.cpu().detach().numpy() / np.linalg.
↳norm(mean_tensor_expert.cpu().detach().numpy())

    # Compute cosine similarities
    similarity_beg = cosine_similarity(cls_test_normalized.reshape(1, -1),
↳cls_normalized_beg.reshape(1, -1))[0][0]
    similarity_exp = cosine_similarity(cls_test_normalized.reshape(1, -1),
↳cls_normalized_exp.reshape(1, -1))[0][0]

    # Classify based on similarities
    predicted = "beginner" if similarity_beg > similarity_exp else "expert"

    # Update confusion matrix counters based on prediction and truth
    if predicted == "beginner":
```

```

        if truth == "beginner":
            tp_beg += 1
        else:
            fp_beg += 1
    else: # predicted == "expert"
        if truth == "expert":
            tp_exp += 1
        else:
            fp_exp += 1

    if predicted != truth: # False negatives
        if truth == "beginner":
            fn_beg += 1
        else:
            fn_exp += 1

tp_total = tp_beg + tp_exp
fp_total = fp_beg + fp_exp
fn_total = fn_beg + fn_exp

# Calculate precision, recall, and F1 score for overall classification
precision_total = tp_total / (tp_total + fp_total) if (tp_total + fp_total) > 0
    else 0
recall_total = tp_total / (tp_total + fn_total) if (tp_total + fn_total) > 0
    else 0
f1_total = 2 * (precision_total * recall_total) / (precision_total +
    recall_total) if (precision_total + recall_total) > 0 else 0

# Print the overall F1 score
print(f"Overall F1 score: {f1_total:.4f}")

# Create confusion matrix
conf_matrix = np.array([[tp_beg, fp_exp],
                        [fn_beg, tp_exp]])

# Plot the confusion matrix using seaborn's heatmap
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
    xticklabels=["Beginner", "Expert"], yticklabels=["Beginner", "Expert"],
    cbar=False)

# Set plot labels and title
plt.title("Confusion Matrix", fontsize=16)
plt.xlabel("Predicted Label", fontsize=12)
plt.ylabel("True Label", fontsize=12)

# Display the plot

```

```
plt.show()
```

Overall F1 score: 0.9608

